

# Final Review 02



Operating Systems  
Wenbo Shen

# Summary

---

- Computer architecture
- OS introduction
- OS structures
- Processes
- IPC
- Thread
- Scheduling
- Synchronization
- Deadlock

# Summary

---

- Memory – segmentation
- Memory – paging
- Virtual memory
- Virtual memory – Linux
- Mass storage
- IO
- FS interface
- FS implementation
- FS in practice

# 04: Thread



# Motivation

---

- Why threads?
  - multiple tasks of an application can be implemented by threads
    - e.g., update display, fetch data, spell checking, answer a network request
  - process creation is heavy-weight while thread creation is light-weight - why?
  - threads can simplify code, increase efficiency
- Kernels are generally multithreaded

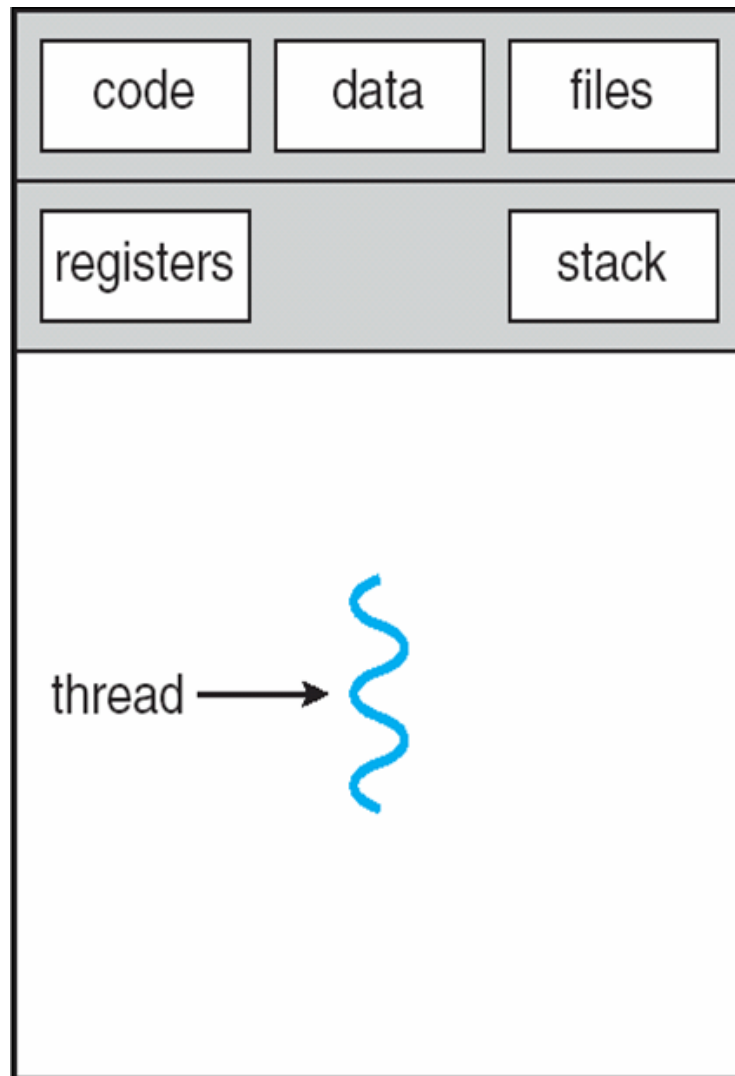
# Thread Definition

---

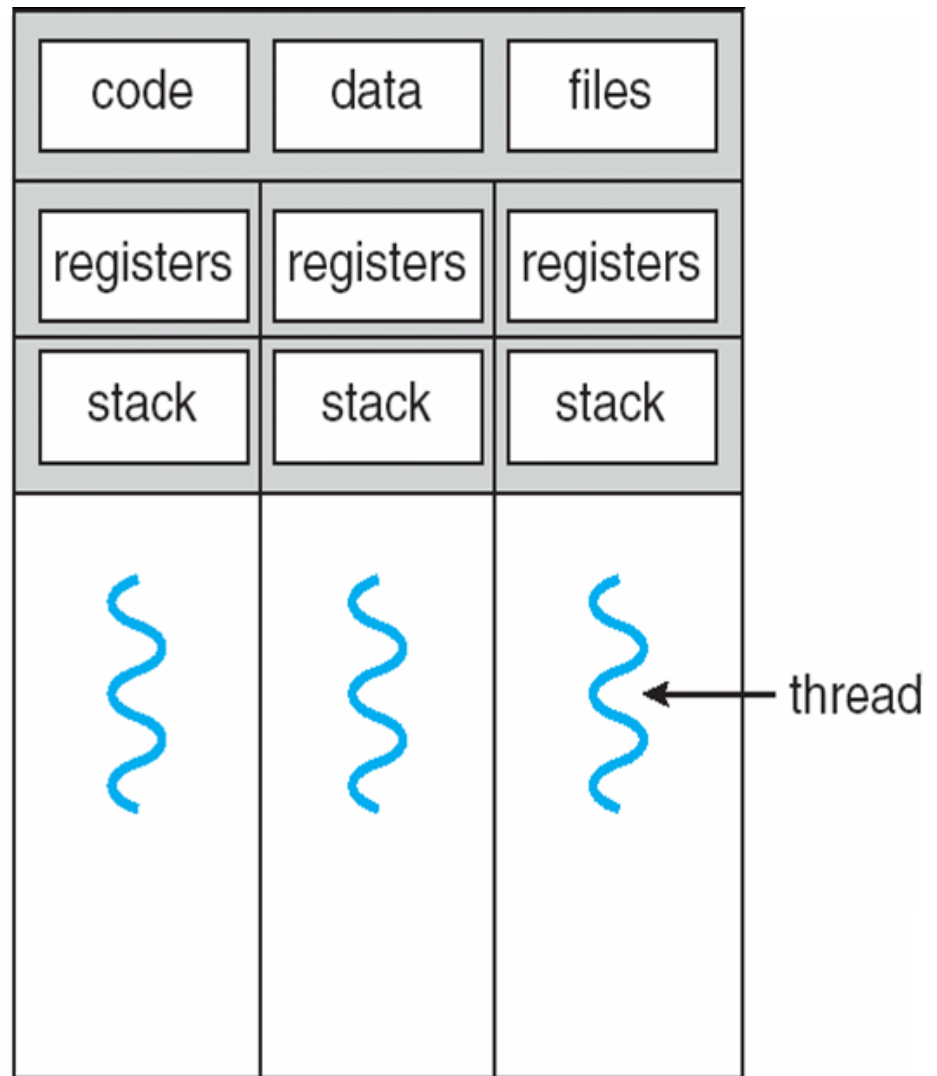
- A thread is a basic unit of execution within a process
- Each thread has its own
  - thread ID
  - program counter
  - register set
  - Stack
- It shares the following with other threads within the same process
  - code section
  - data section
  - the heap (dynamically allocated memory)
  - open files and signals
- **Concurrency:** A multi-threaded process can do multiple things at once

# The Typical Figure

---

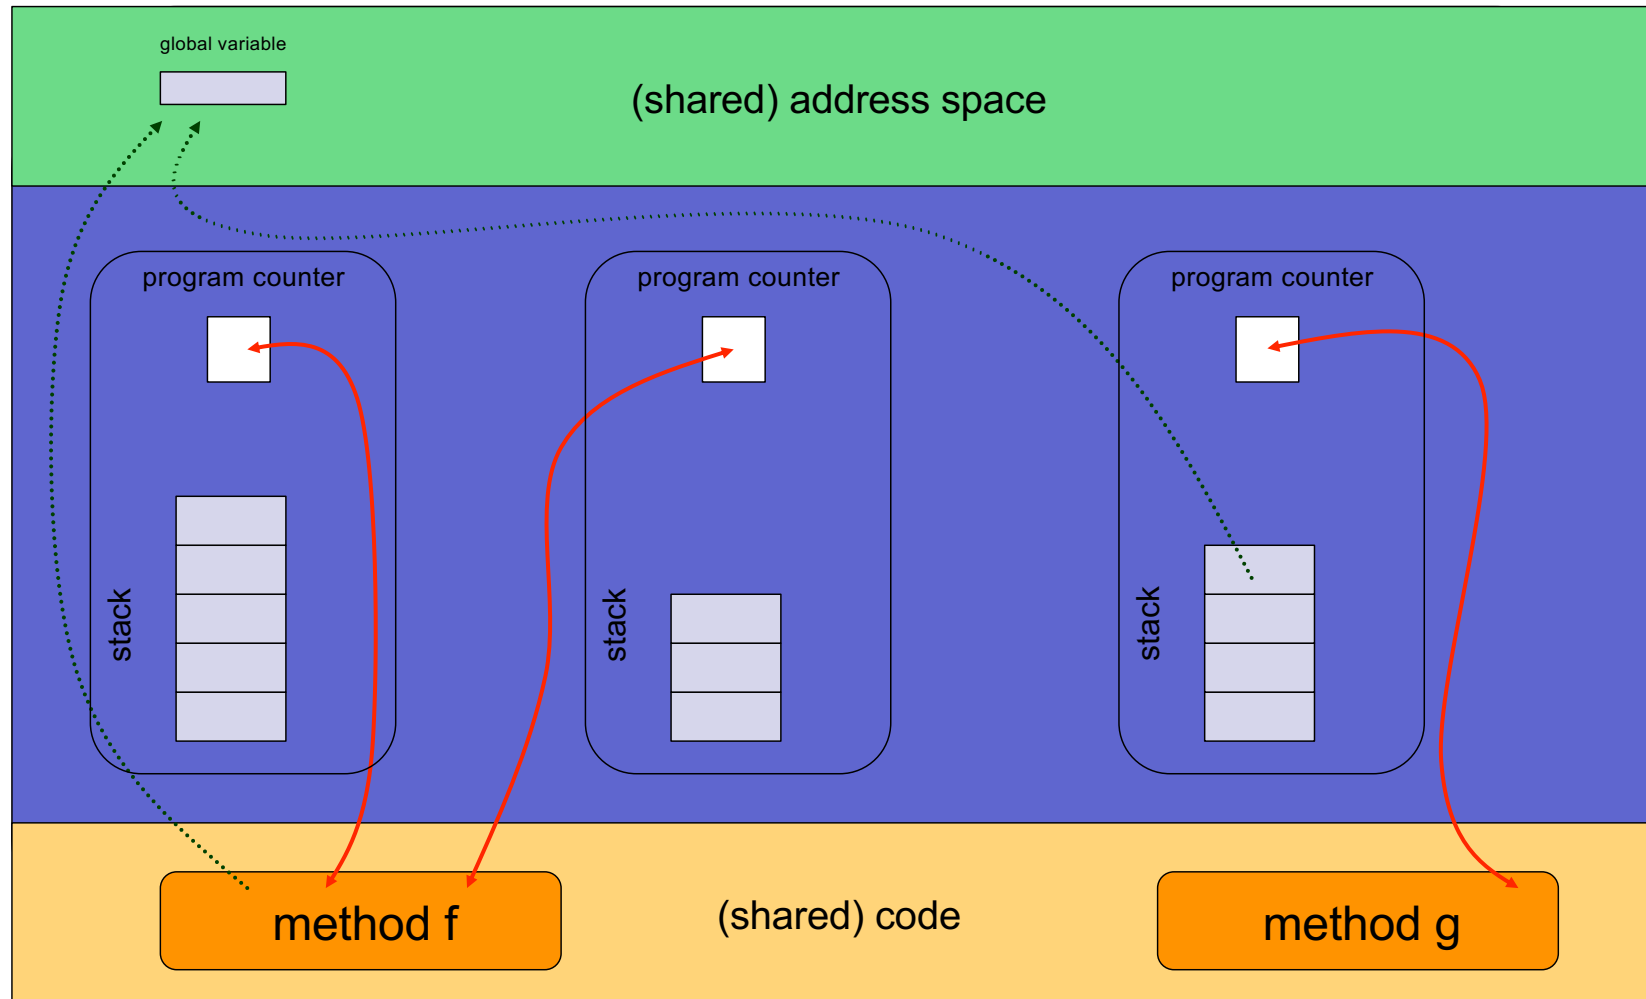


single-threaded process



multithreaded process

# Thread and Process





# Advantages of Threads

---

- Economy:
  - Creating a thread is cheap
    - Much cheaper than creating a process
      - Code, data and heap are already in memory
  - Context-switching between threads is cheap
    - Much cheaper than between processes
      - No cache flush
- Resource Sharing:
  - Threads naturally share memory
    - With processes you have to use possibly complicated IPC (e.g., Shared Memory Segments)
    - **IPC is not needed**
  - Having concurrent activities in the same address space is very powerful
    - But fraught with danger

# Advantages of Threads?

---

- Responsiveness
  - A program that has concurrent activities is more responsive
    - While one thread blocks waiting for some event, another can do something
    - e.g. Spawn a thread to answer a client request in a client-server implementation
  - This is true of processes as well, but with threads we have better sharing and economy
- Scalability
  - Running multiple “threads” at once uses the machine more effectively
    - e.g., on a multi-core machine
  - This is true of processes as well, but with threads we have better sharing and economy

# Drawbacks of Threads

---

- Weak isolation between threads: If **one thread fails** (e.g., a segfault), then **the process fails**
  - And therefore the whole program
- Threads may be more memory-constrained than processes
  - Due to OS limitation of the address space size of a single process
  - Not a problem any more on 64-bit architecture
- Threads do not benefit from memory protection
  - Concurrent programming with Threads is hard
    - But so is it with Processes and Shared Memory Segments

# Implementing Threads

---

- Thread may be provided either at the user level, or by the kernel
  - User threads are supported above the kernel and managed without kernel support
    - Three thread libraries: POSIX pthreads, win32 threads, and java threads
  - Kernel threads are supported and managed directly by the kernel
    - All contemporary OS supports kernel threads

# Multithreading Models

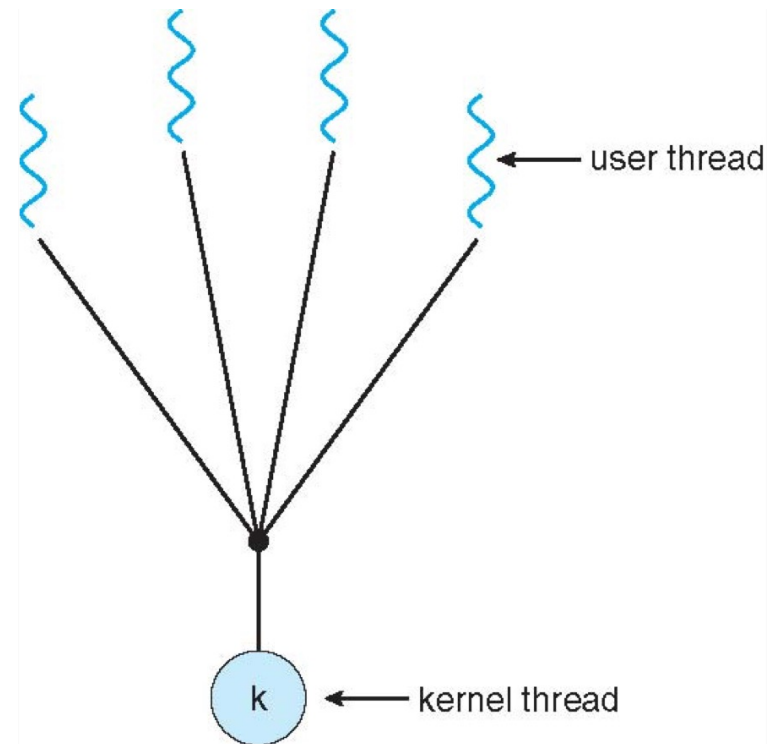
---

- A relationship **must exist** between user threads and kernel threads
  - Kernel threads are the real threads in the system, so for a user thread to make progress the user program has to have its scheduler take a user thread and then run it on a kernel thread.

# Many-to-One

---

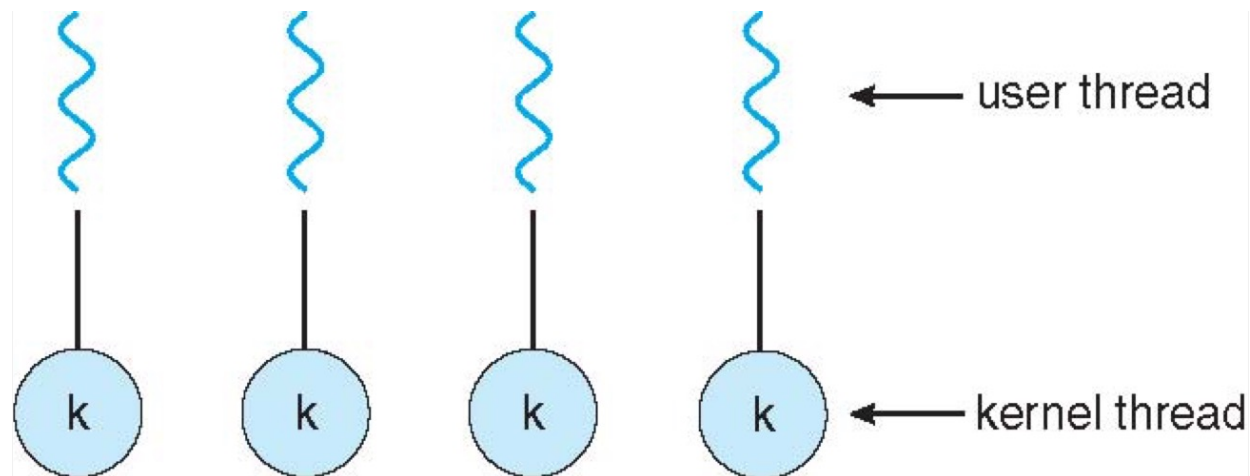
- Many user-level threads mapped to a single kernel thread
  - Thread management is done by the thread library in **user space** (efficient)
  - Entire process will block if a thread makes a blocking system call
    - Convert blocking system call to non-blocking (e.G., Select in unix)?
  - Multiple threads are unable to run in parallel on multi-processors
- Examples:
  - Solaris green threads



# One-to-One

---

- Each user-level thread maps to one kernel thread
  - It allows other threads to run when a thread blocks
  - Multiple thread can run in parallel on multiprocessors
  - Creating a user thread requires creating a corresponding kernel thread
    - It leads to overhead
  - Most operating systems implementing this model limit the number of threads
- Examples
  - Windows NT/XP/2000
  - Linux



# Many-to-Many Model

---

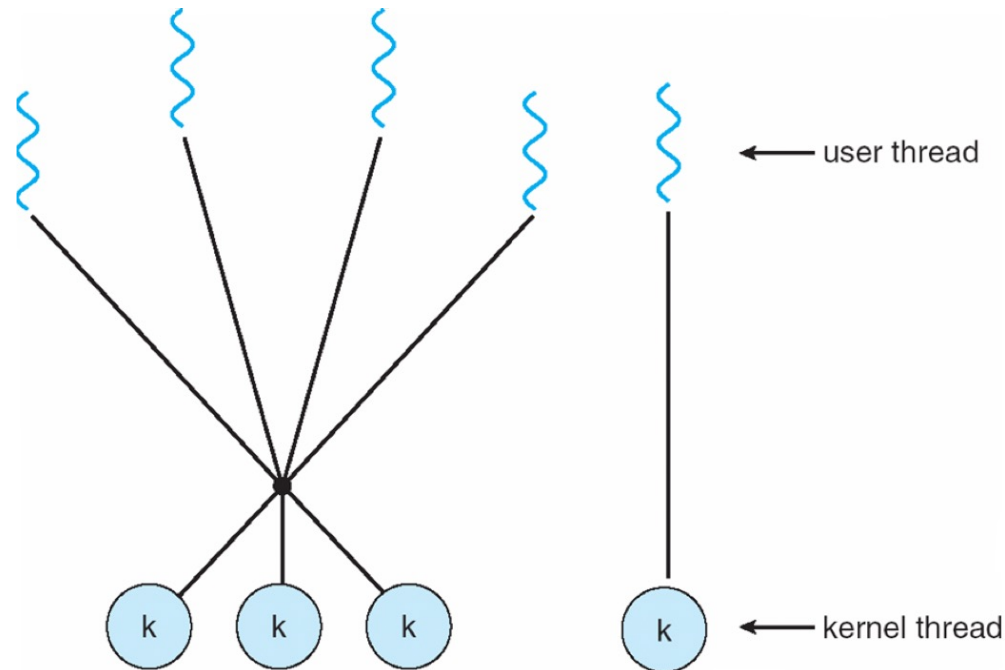
- Many user level threads are mapped to many kernel threads
  - it solves the shortcomings of 1:1 and m:1 model
  - developers can create as many user threads as necessary
  - corresponding kernel threads can run in parallel on a multiprocessor
- Examples
  - Solaris prior to version 9
  - Windows NT/2000 with the ThreadFiber package



# Two-level Model

---

- Similar to many-to-many model, except that it allows a user thread to be **bound** to kernel thread



# Semantics of Fork and Exec

---

- Fork duplicates the whole single-threaded process
- Does fork duplicate only the calling thread or all threads for multi-threaded process?
  - some UNIX systems have two versions of fork, one for each semantic
- Exec typically replaces the entire process, multithreaded or not
  - use “fork the calling thread” if calling exec soon after fork
- Which version of fork to use depends on the application
  - Exec is called immediately after forking: duplicating all threads is not necessary
  - Exec is not called: duplicating all threads

# Linux Threads

---

- Linux does not distinguish between PCB and TCB
  - Kernel data structure: task\_struct

```
591
592 struct task_struct {
593 #ifdef CONFIG_THREAD_INFO_IN_TASK
594     /*
595      * For reasons of header soup (see current_thread_info()), this
596      * must be the first element of task_struct.
597      */
598     struct thread_info          thread_info;
599 #endif
600     /* -1 unrunnable, 0 runnable, >0 stopped: */
601     volatile long               state;
602
603     /*
604      * This begins the randomizable portion of task_struct. Only
605      * scheduling-critical items should be added above here.
606      */
607     randomized_struct_fields_start
608
609     void                        *stack;
610     atomic_t                    usage;
611     /* Per task flags (PF_*), defined further below: */
612     unsigned int                flags;
613     unsigned int                ptrace;
614
```

# Linux Threads

---

- In Linux, a thread is also called a light-weight process (LWP)
- The clone() syscall is used to create a thread or a process
  - Shares execution context with its parent
  - pthread library uses clone() to implement threads. Refer to `./nptl/sysdeps/pthread/createthread.c`

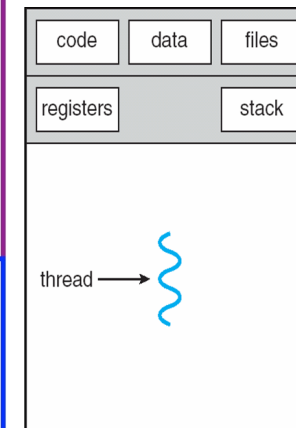
flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

# Linux Threads

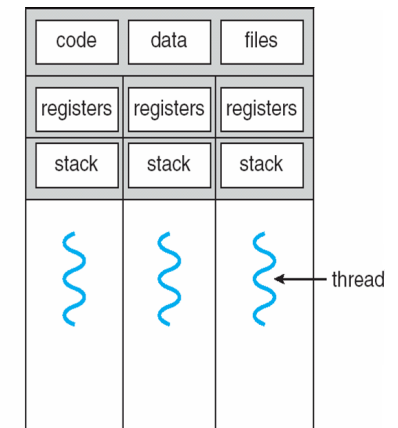
- Single-threaded process vs multi-threaded process

```
wenbo@wenbo-desktop:~/KERNEL/linux.git$ ps -eLf
```

UID	PID	PPID	LWP	C	NLWP	STIME	TTY	TIME	CMD
root	1	0	1	0	1	3月11 ?		00:00:19	/sbin/init splash
root	2	0	2	0	1	3月11 ?		00:00:00	[kthreadd]
root	4	2	4	0	1	3月11 ?		00:00:00	[kworker/0:0H]
root	6	2	6	0	1	3月11 ?		00:00:00	[mm_percpu_wq]
root	7	2	7	0	1	3月11 ?		00:00:00	[ksoftirqd/0]
root	8	2	8	0	1	3月11 ?		00:00:31	[rcu_sched]
root	9	2	9	0	1	3月11 ?		00:00:00	[rcu_bh]
root	10	2	10	0	1	3月11 ?		00:00:00	[migration/0]
root	11	2	11	0	1	3月11 ?		00:00:00	[watchdog/0]
root	704	1	704	0	1	3月11 ?		00:00:00	/usr/sbin/cron -f
root	718	1	718	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	882	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	883	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	884	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	885	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	917	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	921	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd
root	718	1	922	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	923	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd
root	718	1	924	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd



single-threaded process



multithreaded process

# Linux Threads

- Single-threaded process vs multi-threaded process

```
wenbo@wenbo-desktop:~/KERNEL/linux.git$ ps -eLf
```

UID	PID	PPID	LWP	C	NLWP	STIME	TTY	TIME	CMD
root	1	0	1	0	1	3月11 ?		00:00:19	/sbin/init splash
root	2	0	2	0	1	3月11 ?		00:00:00	[kthreadd]
root	4	2	4	0	1	3月11 ?		00:00:00	[kworker/0:0H]
root	6	2	6	0	1	3月11 ?		00:00:00	[mm_percpu_wq]
root	7	2	7	0	1	3月11 ?		00:00:00	[ksoftirqd/0]
root	8	2	8	0	1	3月11 ?		00:00:31	[rcu_sched]
root	9	2	9	0	1	3月11 ?		00:00:00	[rcu_bh]
root	10	2	10	0	1	3月11 ?		00:00:00	[migration/0]
root	11	2	11	0	1	3月11 ?		00:00:00	[watchdog/0]
root	704	1	704	0	1	3月11 ?		00:00:00	/usr/sbin/cron -f
root	718	1	718	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	882	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	883	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	884	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	885	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	917	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	921	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd
root	718	1	922	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	923	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd
root	718	1	924	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd

```
787      /* PID/PID hash table linkage. */
788      struct pid                *thread_pid;
789      struct hlist_node         pid_links[PIDTYPE];
790      struct list_head          thread_group;
791      struct list_head          thread_node;
792
793      struct completion          *vfork_done;
794
795      /* CLONE_CHILD_SETTID: */
796      int __user                 *set_child_tid;
797      ---
```



# Threads with Process – What is shared

---

```
--
29 static void traversal_thread_group(struct task_struct * tsk){
30     struct task_struct * curr_thread = NULL;
31     unsigned long tg_offset = offsetof(struct task_struct, thread_group);
32
33     curr_thread = (struct task_struct *) (((unsigned long)tsk->thread_group.next) - tg_offset);
34     while (curr_thread != tsk){
35         printk("\t\tTHREAD TSK=%llx\tPID=%d\tSTACK=%llx \tCOMM=%s\tMM=%llx\tACTIVE_MM=%llx\n",
36             (u64)curr_thread, curr_thread->pid, (u64)curr_thread->stack,
37             curr_thread->comm, (u64)curr_thread->mm, (u64)curr_thread->active_mm);
38         curr_thread = (struct task_struct *) (((unsigned long)curr_thread->thread_group.next) - tg_offset);
39     }
40 }
41
42 static void traversal_process(void) {
43     struct task_struct * tsk = NULL;
44
45     traversal_thread_group(&init_task);
46     for_each_process(tsk){
47         printk("PROCESS\tTHREAD TSK=%llx\tPID=%d\tSTACK=%llx \tCOMM=%s\tMM=%llx\tACTIVE_MM=%llx\n",
48             (u64)tsk, tsk->pid, (u64)tsk->stack, tsk->comm,
49             (u64)tsk->mm, (u64)tsk->active_mm);
50         traversal_thread_group(tsk);
51     }
52 }
```

# Threads with Process – What is shared

```

29 static void traversal_thread_group(struct task_struct * tsk){
30     struct task_struct * curr_thread = NULL;
31     unsigned long tg_offset = offsetof(struct task_struct, thread_group);
32
33     curr_thread = (struct task_struct *) (((unsigned long)tsk->thread_group.next) - tg_offset);
34     while (curr_thread != tsk){
35         printk("\t\tTHREAD TSK=%llx\tPID=%d\tSTACK=%llx \tCOMM=%s\tMM=%llx\tACTIVE_MM=%llx\n",
36             (u64)curr_thread, curr_thread->pid, (u64)curr_thread->stack,
37             curr_thread->comm, (u64)curr_thread->mm, (u64)curr_thread->active_mm);
38         curr_thread = (struct task_struct *) (((unsigned long)curr_thread->thread_group.next) - tg_offset);
39     }
40 }
41
42 static void traversal_process(void) {
43     struct task_struct * tsk = NULL;
44
45     traversal_thread_group(&init_task);
46     for_each_process(tsk){
47         printk("PROCESS\t\tTHREAD TSK=%llx\tPID=%d\tSTACK=%llx \tCOMM=%s\tMM=%llx\tACTIVE_MM=%llx\n",
48             (u64)tsk, tsk->pid, (u64)tsk->stack, tsk->comm,
49             (u64)tsk->mm, (u64)tsk->active_mm);
50         traversal_thread_group(tsk);
51     }
52 }

```

PROCESS	THREAD	TSK=ffff8c4c4bf3c5c0	PID=718	STACK=ffff985c82268000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46d52e80	PID=882	STACK=ffff985c82390000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46d545c0	PID=883	STACK=ffff985c822e8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c491b45c0	PID=884	STACK=ffff985c8218c000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4c4beb1740	PID=885	STACK=ffff985c821ec000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4c4ae1ae80	PID=917	STACK=ffff985c823c8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4c4b562e80	PID=921	STACK=ffff985c82418000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4c48340000	PID=922	STACK=ffff985c823b0000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4c472bae80	PID=923	STACK=ffff985c821f4000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4c4b5945c0	PID=924	STACK=ffff985c81fa8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4c46775d00	PID=925	STACK=ffff985c822a8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4c4b692e80	PID=973	STACK=ffff985c82438000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4c4b78ae80	PID=974	STACK=ffff985c823c0000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4c46e1dd00	PID=975	STACK=ffff985c824b8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840



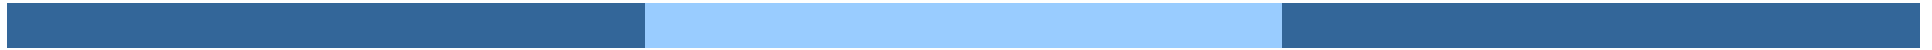
# Threads within Process – What is shared

PROCESS	THREAD	TSK=ffff8c4c4bf3c5c0	PID=718	STACK=ffff985c82268000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46d52e80	PID=882	STACK=ffff985c82390000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46d545c0	PID=883	STACK=ffff985c822e8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c491b45c0	PID=884	STACK=ffff985c8218c000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4be1740	PID=885	STACK=ffff985c821ec000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4ae1ae80	PID=917	STACK=ffff985c823c8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b562e80	PID=921	STACK=ffff985c82418000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c48340000	PID=922	STACK=ffff985c823b0000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c472bae80	PID=923	STACK=ffff985c821f4000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b5945c0	PID=924	STACK=ffff985c81fa8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46775d00	PID=925	STACK=ffff985c822a8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b692e80	PID=973	STACK=ffff985c82438000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b78ae80	PID=974	STACK=ffff985c823c0000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46e1dd00	PID=975	STACK=ffff985c824b8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
		task_struct	pid	stack	comm		mm_struct

Not  
Shared

Shar  
ed

# 05: CPU Scheduling



# Basic Concepts

---

- Process execution consists of a cycle of CPU execution and I/O wait
  - CPU burst and I/O burst alternate
  - CPU burst distribution varies greatly from process to process, and from computer to computer, but follows similar curves
  - **Rationale**: non-CPU-intensive jobs should really get the CPU quickly on the rare occasions they need them, because they could be interactive processes
  - Maximum CPU utilization obtained with multiprogramming

# CPU Scheduler

---

- CPU scheduler selects from among the processes in **ready queue**, and allocates the CPU to one of them
- CPU scheduling decisions **may take place** when a process:
  - Switches from **running to waiting state** (e.G., Wait for I/O)
  - Switches from **running to ready state** (e.G., When an interrupt occurs)
  - Switches from **waiting to ready** (e.G., At completion of I/O)
  - **Terminates**
- Scheduling under condition **1 and 4 only** is **nonpreemptive**
  - Once the CPU has been allocated to a process, the process keeps it until terminates or waiting for I/O
  - Also called **cooperative scheduling**
- **Preemptive scheduling** schedules process **also** in condition **2 and 3**
  - Preemptive scheduling needs hardware support such as a timer
  - Synchronization primitives are necessary
- Context switch can only happen in kernel mode, so is preemption
  - User space processes need to trap to kernel mode to do context switch.

# Scheduling Criteria

---

- **CPU utilization** : percentage of CPU being busy
- Throughput: # of processes that complete execution per time unit
- Turnaround time: the time to execute a particular process
  - From the time of *submission* to the time of *completion*
- **Waiting time**: the total time spent waiting in the *ready queue*
- Response time: the time it takes from when a request was submitted until the first response is produced
  - The time it takes to *start responding*

# Scheduling Algorithms

---

- First-come, first-served scheduling (FCFS)
- Shortest-job-first scheduling (SJF)
- Priority scheduling
- Round-robin scheduling (RR)
- Multilevel queue scheduling
- Multilevel feedback queue scheduling

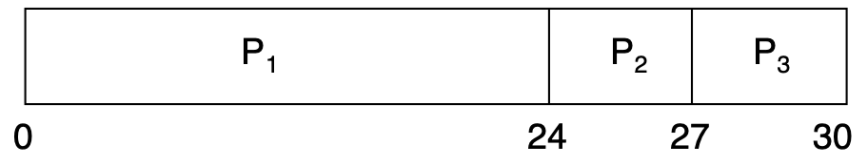
# First-Come, First-Served (FCFS) Scheduling

---

- Example processes:

Process	Burst Time
P <sub>1</sub>	24
P <sub>2</sub>	3
P <sub>3</sub>	3

- Suppose that the processes arrive in the order: P<sub>1</sub> , P<sub>2</sub> , P<sub>3</sub>
- the Gantt Chart for the FCFS schedule is:



- **Waiting time** for P<sub>1</sub> = 0; P<sub>2</sub> = 24; P<sub>3</sub> = 27, **average waiting time**:  $(0 + 24 + 27)/3 = 17$

# Shortest-Job-First Scheduling

---

- Associate with each process: the length of its next CPU burst
  - the process with the **smallest next CPU burst** is scheduled to run next
- SJF is **provably optimal**: it gives **minimum average waiting** time for a given set of processes
  - moving a short process before a long one decreases the overall waiting time
  - the difficulty is to know the length of the next CPU request
    - long-term scheduler can use the user-provided processing time estimate
    - short-term scheduler needs to approximate SFJ scheduling
- SJF can be **preemptive** or **nonpreemptive**
  - preemptive version is called **shortest-remaining-time-first**

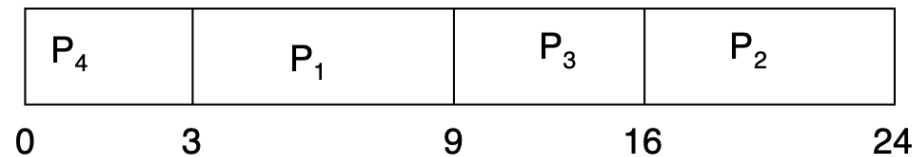


# Example of SJF

---

Process	Burst Time
$P_1$	6
$P_2$	8
$P_3$	7
$P_4$	3

- SJF scheduling chart



- **Average waiting time =  $(3 + 16 + 9 + 0) / 4 = 7$**

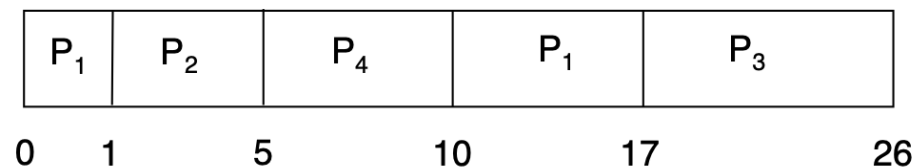
# Shortest-Remaining-Time-First

---

- SJF can be **preemptive: reschedule when a process arrives**

Process	Arrival Time	Burst Time
P <sub>1</sub>	0	8
P <sub>2</sub>	1	4
P <sub>3</sub>	2	9
P <sub>4</sub>	3	5

- Preemptive SJF Gantt Chart



- Average waiting time =  $[(10-1)+(1-1)+(17-2)+5-3]/4 = 26/4 = 6.5$  msec

# Priority Scheduling

---

- Priority scheduling selects the ready process with **highest priority**
  - a priority number is associated with each process, smaller integer, higher priority
  - the CPU is allocated to the process with the highest priority
  - SJF is special case of priority scheduling
    - priority is the inverse of predicted next CPU burst time
- Priority scheduling can be **preemptive** or **nonpreemptive**, similar to SJF
- **Starvation** is a problem: **low priority processes may never execute**
  - **Solution: aging** — gradually increase priority of processes that wait for a long time

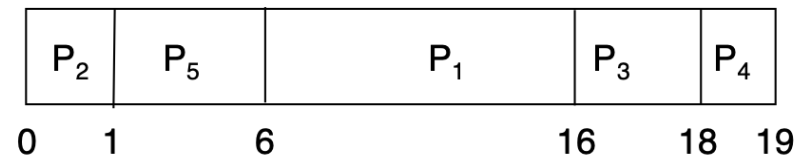
# Example of Priority Scheduling

---

Process Burst Time Priority

P <sub>1</sub>	10	3
P <sub>2</sub>	1	1
P <sub>3</sub>	2	4
P <sub>4</sub>	1	5
P <sub>5</sub>	5	2

- Priority scheduling Gantt Chart



- Average waiting time = 8.2 msec

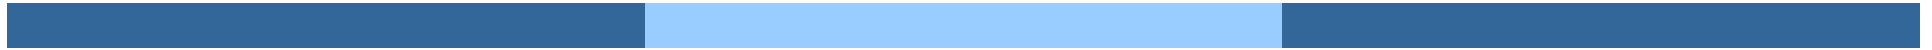
We use small number to denote high priority.

# Round Robin (RR)

---

- Round-robin scheduling selects process in a **round-robin** fashion
  - each process gets a small unit of CPU time (time quantum,  $q$ )
    - $q$  is too large  $\rightarrow$  FIFO,  $q$  is too small  $\rightarrow$  context switch overhead is high
    - a time quantum is generally 10 to 100 milliseconds

# 06&07: Synchronization



# Background

---

- Processes can execute concurrently
  - May be interrupted at any time, partially completing execution
- Concurrent access to shared data may result in **data inconsistency**
  - data consistency requires orderly execution of cooperating processes

# Uncontrolled Scheduling

- Counter = counter + 1

```
mov 0x8049a1c, %eax
add $0x1, %eax
mov %eax, 0x8049a1c
```

OS	Thread 1	Thread 2	(after instruction)		
			PC	%eax	counter
	<i>before critical section</i>		100	0	50
	mov 0x8049a1c, %eax		105	<b>50</b>	50
	add \$0x1, %eax		108	<b>51</b>	50
<b>interrupt</b>	<i>save T1's state</i>				
	<i>restore T2's state</i>		100	0	50
		mov 0x8049a1c, %eax	105	<b>50</b>	50
		add \$0x1, %eax	108	<b>51</b>	50
		mov %eax, 0x8049a1c	113	51	<b>51</b>
<b>interrupt</b>	<i>save T2's state</i>				
	<i>restore T1's state</i>		108	51	51
	mov %eax, 0x8049a1c		113	51	<b>51</b>

**counter: 51 instead of 52!**



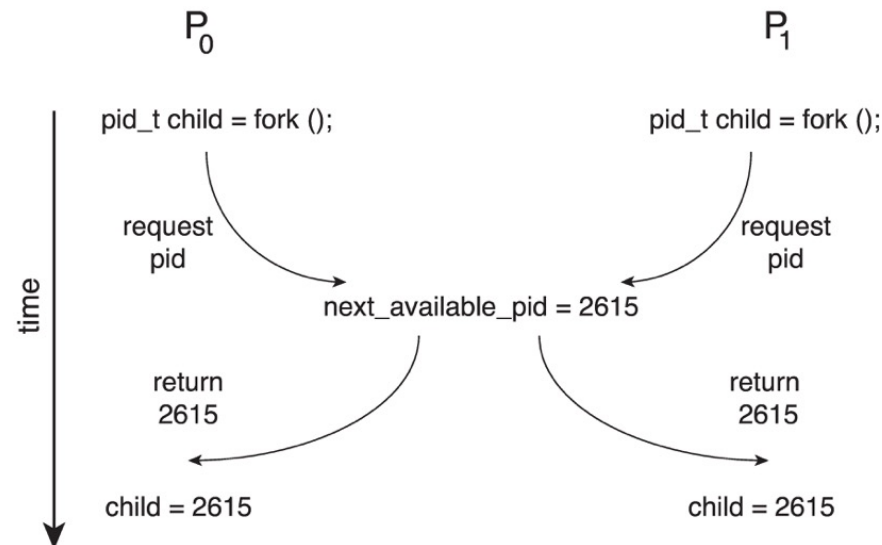
# Race Condition

---

- Several processes (or threads) access and manipulate the same data concurrently and the outcome of the execution depends on the particular order in which the access takes place, is called a **race-condition**

# Race Condition in Kernel

- Processes P0 and P1 are creating child processes using the fork() system call
- Race condition on kernel variable ***next\_available\_pid*** which represents the next available process identifier (pid)



- Unless there is mutual exclusion, the same pid could be assigned to two different processes!
- Even if the kernel is non-preemptive, race condition can still exist in user space!**

# Critical Section

---

- General structure of process  $p_i$  is

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```

# Critical-Section Handling in OS

---

- Single-core system: preventing interrupts
- Multiple-processor: preventing interrupts are not feasible
- Two approaches depending on if kernel is ***preemptive or non-preemptive***
  - Preemptive – allows preemption of process when running in kernel mode
  - Non-preemptive – runs until **exits kernel mode, blocks, or voluntarily yields CPU**
    - Essentially free of race conditions ***in kernel mode, but NOT for user space!!***

## Solution to Critical-Section: Three Requirements

---

- **Mutual Exclusion**
  - only one process can execute in the critical section
- **Progress**
- **Bounded waiting**
  - it prevents **starvation**

# Peterson's Solution

---

- Peterson's solution solves **two-processes** synchronization
- **It's a software based-solution**
- It assumes that LOAD and STORE are **atomic**
  - **atomic**: execution cannot be interrupted
- The two processes share two variables
  - int **turn**: whose turn it is to enter the critical section
  - Boolean **flag[2]**: whether a process is ready to enter the critical section

# Peterson's Solution

---

```
flag[0] = FALSE;  
flag[1] = FALSE;
```

- **P<sub>0</sub>:**

```
do {  
    flag[0] = TRUE;  
    turn = 1;  
    while (flag[1] && turn == 1);  
    critical section  
    flag[0] = FALSE;  
    remainder section  
} while (TRUE);
```

Mark self ready

Assert the other one

- **P<sub>1</sub>:**

```
do {  
    flag[1] = TRUE;  
    turn = 0;  
    while (flag[0] && turn == 0);  
    critical section  
    flag[1] = FALSE;  
    remainder section  
} while (TRUE);
```

# Hardware Instructions

---

- Special hardware instructions that allow us to either test-and-modify the content of a word, or two swap the contents of two words atomically (uninterruptibly.)
- **Test-and-Set** instruction
- **Compare-and-Swap** instruction



# Mutex Locks

---

- OS designers build software tools to solve critical section problem
- Simplest is **mutex lock**
- Protect a critical section by first **acquire()** a lock then **release()** the lock
  - Boolean variable indicating if lock is available or not
- Calls to **acquire()** and **release()** must be **atomic**
  - Usually implemented via hardware atomic instructions such as compare-and-swap.
- But this solution requires **busy waiting**
- This lock therefore called a spinlock

# Mutex Locks

---

```
while (true) {  
    acquire lock  
  
    critical section  
  
    release lock  
  
    remainder section  
}
```

# Mutex Lock Definitions

---

- These two functions must be implemented atomically
  - Both test-and-set and compare-and-swap can be used to implement these functions

```
bool locked = false;
```

```
acquire() {  
    while (compare_and_swap(&locked, false, true))  
        ; //busy waiting  
}
```

```
release() {  
    locked = false;  
}
```

# Semaphore

---

- **Semaphore** S is an integer variable
  - e.g., to represent *how many units of a particular resource is available*
  - *For resource sharing purpose*
- It can only be updated with two atomic operations: **wait** and **signal**
  - **spin lock** can be used to guarantee atomicity of wait and signal
  - originally called P and V (Dutch)
  - a simple implementation with busy wait can be:

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```

# Semaphore

---

- Associate a waiting queue with each semaphore
  - place the process on the waiting queue if **wait** cannot return immediately
  - wake up a process in the waiting queue in **signal**
- There is no need to **busy wait** in critical section
- Note: wait and signal must still be atomic

# Semaphore

---

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}
```

**Suppose the init value s->value = 5  
And now, if s->value = -3, what does it mean?**

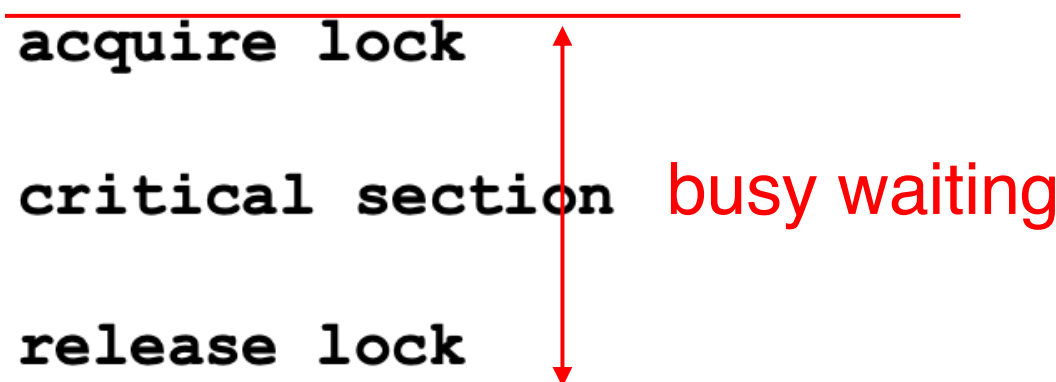
```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a proc.P from S->list;  
        wakeup(P);  
    }  
}
```

# Busy waiting time - Mutex

---

- Mutex busy waiting time
  - From acquire to release

```
while (true) {  
    acquire lock  
  
    critical section busy waiting  
  
    release lock  
  
    remainder section  
}
```

A diagram illustrating the busy waiting period in a mutex. Two horizontal red lines are drawn. The top line is aligned with the 'acquire lock' statement, and the bottom line is aligned with the 'release lock' statement. A vertical red double-headed arrow connects these two lines, spanning the duration of the 'critical section' and the 'remainder section' code blocks. The text 'busy waiting' is written in red to the right of the arrow, indicating the period where the CPU is busy but the lock is not being held.

- What if the critical section is long?
  - A huge waste of CPU time

# Busy waiting time - Semaphore

```
Semaphore sem;    // initialized to 1
do {
    wait (sem);    ⇕ busy waiting
    critical section
    signal (sem);  ⇕ busy waiting
    remainder section
} while (TRUE);    //while loop but not busy waiting
```

- No busy waiting on critical section
- Still has the busy waiting on *wait* and *signal*
  - But waiting is much shorter



# Bounded-Buffer Problem

---

- Two processes, the producer and the consumer share  $n$  buffers
  - the producer generates data, puts it into the buffer
  - the consumer consumes data by removing it from the buffer
- The problem is to make sure:
  - **the producer won't try to add data into the buffer if it is full**
  - **the consumer won't try to remove data from an empty buffer**
  - also call producer-consumer problem

# Bounded-Buffer Problem

---

- Solution:
  - $n$  buffers, each can hold one item
  - semaphore mutex initialized to the value 1
  - semaphore full-slots initialized to the value 0
  - semaphore empty-slots initialized to the value  $N$

# Bounded-Buffer Problem

---

- The producer process:

```
do {  
  //produce an item  
  
  ...  
  
  wait(empty-slots);  
  
  wait(mutex);  
  
  //add the item to the  buffer  
  
  ...  
  
  signal(mutex);  
  
  signal(full-slots);  
  
} while (TRUE)
```

# Bounded Buffer Problem

---

- The consumer process:

```
do {  
    wait(full-slots);  
    wait(mutex);  
    //remove an item from  buffer  
    ...  
    signal(mutex);  
    signal(empty-slots);  
    //consume the item  
    ...  
} while (TRUE);
```

# Takeaway

---

- Whole slides