

Final Review 03



Operating Systems
Wenbo Shen


Summary

- Computer architecture
- OS introduction
- OS structures
- Processes
- IPC
- Thread
- Scheduling
- Synchronization
- Deadlock

Summary

- Memory – segmentation
- Memory – paging
- Virtual memory
- Virtual memory – Linux
- Mass storage
- IO
- FS interface
- FS implementation
- FS in practice

08: Deadlock



The Deadlock Problem

- **Deadlock:** a set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set
- Examples:
 - a system has 2 disk drives, P_1 and P_2 each hold one disk drive and each needs another one
 - semaphores A and B, initialized to 1

P_1	P_2
wait (A);	wait(B)
wait (B);	wait(A)

Deadlock in program


- Two mutex locks are created and initialized:

```
pthread_mutex_t first_mutex;  
pthread_mutex_t second_mutex;  
  
pthread_mutex_init(&first_mutex, NULL);  
pthread_mutex_init(&second_mutex, NULL);
```

```
/* thread one runs in this function */  
void *do_work_one(void *param){  
    pthread_mutex_lock(&first_mutex);  
    pthread_mutex_lock(&second_mutex);  
  
    /* Do some work*/  
  
    pthread_mutex_unlock(&second_mutex);  
    pthread_mutex_unlock(&first_mutex);  
    pthread_exit(0);  
}  
  
/* thread two runs in this function */  
void *do_work_two(void *param){  
    pthread_mutex_lock(&second_mutex);  
    pthread_mutex_lock(&first_mutex);  
  
    /* Do some work*/  
  
    pthread_mutex_unlock(&first_mutex);  
    pthread_mutex_unlock(&second_mutex);  
  
    pthread_exit(0);  
}
```

Deadlock in program

- Deadlock is possible if thread 1 acquires **first_mutex** and thread 2 acquires **second_mutex**. Thread 1 then waits for **second_mutex** and thread 2 waits for **first_mutex**.
- Can be illustrated with a **resource allocation graph**:



Four Conditions of Deadlock

- **Mutual exclusion:** only one process at a time can use a resource
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
- **No preemption:** a resource can be released only voluntarily by the process holding it, after it has completed its task
- **Circular wait:** there exists a set of waiting processes $\{P_0, P_1, \dots, P_n\}$
 - P_0 is waiting for a resource that is held by P_1
 - P_1 is waiting for a resource that is held by P_2 ...
 - P_{n-1} is waiting for a resource that is held by P_n
 - P_n is waiting for a resource that is held by P_0

How to Handle Deadlocks

- Ensure that the system will never enter a deadlock state
 - **Prevention**
 - **Avoidance**
- Allow the system to enter a deadlock state and then recover - database
 - **Deadlock detection and recovery:**
- **Ignore the problem** and pretend deadlocks never occur in the system
 - Usually adopted by **mainstream operating systems**



Deadlock Prevention

- How to prevent **mutual exclusion**
 - Not required for sharable resources
 - Must hold for non-sharable resources
- How to prevent **hold and wait**
 - Whenever a process requests a resource, it doesn't hold any other resources
 - Require process to request ***all*** its resources before it begins execution
 - Allow process to request resources only when the process has none
 - Low resource utilization; starvation possible

Deadlock Prevention

- How to handle **no preemption**
 - if a process requests a resource not available
 - release all resources currently being held
 - preempted resources are added to the list of resources it waits for
 - process will be restarted only when it can get all waiting resources
- How to handle **circular wait**
 - impose a total ordering of all resource types
 - require that each process requests resources in an increasing order
 - Many operating systems adopt this strategy for some locks.

Deadlock Avoidance

- **Dead avoidance: require extra information about how resources are to be requested**
 - **Is this requirement practical?**
- Each process declares a **max** number of resources it may need
- Deadlock-avoidance algorithm ensure there can **never be a circular-wait condition**
- Resource-allocation state:
 - the number of **available** and **allocated** resources
 - the **maximum demands** of the processes

Deadlock Avoidance Algorithms

- Single instance of each resource type \Rightarrow use **resource-allocation graph**
- Multiple instances of a resource type \Rightarrow use the **banker's algorithm**
- Allocate based on safe state
 - Limit the order of resource application

Banker's Algorithm

- Banker's algorithm is for **multiple-instance resource deadlock avoidance**
 - each process must claim **maximum** use of each resource type **in advance**
 - when a process requests a resource it may have to wait
 - when a process gets all its resources it must release them in a finite amount of time

Data Structures for the Banker's Algorithm

- n processes, m types of resources
 - available: an array of length m , instances of available resource
 - $\text{available}[j] = k$: k instances of resource type R_j available
 - max: a $n \times m$ matrix
 - $\text{max}[i,j] = k$: process P_i may request at most k instances of resource R_j
 - allocation: $n \times m$ matrix
 - $\text{allocation}[i,j] = k$: P_i is currently allocated k instances of R_j
 - need: $n \times m$ matrix
 - $\text{need}[i,j] = k$: P_i may need k more instances of R_j to complete its task
 - **$\text{need}[i,j] = \text{max}[i,j] - \text{allocation}[i,j]$**

Banker's Algorithm: Example

- System state:
 - **5 processes** P_0 through P_4
 - **3 resource types**: A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time T_0 :

	allocation	max	available
	A B C	A B C	A B C
P_0	0 1 0	7 5 3	3 3 2
P_1	2 0 0	3 2 2	
P_2	3 0 2	9 0 2	
P_3	2 1 1	2 2 2	
P_4	0 0 2	4 3 3	

Banker's Algorithm: Safe State

- Data structure to compute whether the system is in a safe state
 - use **work** (a vector of length m) to track **allocable resources**
 - **current available resources**
 - **unallocated + released by finished processes**
 - use **finish** (a vector of length n) to track whether process has finished
 - initialize: **work** = **available**, **finish[i] = false** for $i = 0, 1, \dots, n-1$
- Algorithm:
 - find an i such that **finish[i] = false** && **need[i] ≤ work**
 - if no such i exists, go to step 3
 - **work = work + allocation[i]**, **finish[i] = true**, go to step 1
 - if **finish[i] == true** for all i , then the system is in a safe state

Bank's Algorithm: Resource Allocation

- Data structure: request vector for process P_i
 - **request**[j] = k then process P_i wants k instances of resource type R_j
- Algorithm:
 1. if $\text{request}[i] \leq \text{need}[i]$ go to step 2; otherwise, raise error condition (the process has exceeded its maximum claim)
 2. if $\text{request}[i] \leq \text{available}$, go to step 3; otherwise P_i must wait (not all resources are not available)
 3. pretend to allocate requested resources to P_i by modifying the state:
 - $\text{available} = \text{available} - \text{request}[i]$
 - $\text{allocation}[i] = \text{allocation}[i] + \text{request}[i]$
 - $\text{need}[i] = \text{need}[i] - \text{request}[i]$
 4. use **previous algorithm** to test if it is a safe state, if so \Rightarrow allocate the resources to P_i
 5. if unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored

Banker's Algorithm: Example

- System state:
 - **5 processes** P_0 through P_4
 - **3 resource types**: A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time T_0 :

	allocation			max		
	A	B	C	A	B	C
P_0	0	1	0	7	5	3
P_1	2	0	0	3	2	2
P_2	3	0	2	9	0	2
P_3	2	1	1	2	2	2
P_4	0	0	2	4	3	3

available		
A	B	C
3	3	2

Banker's Algorithm: Example

- System state:
 - **5 processes** P_0 through P_4
 - **3 resource types**: A (10 instances), B (5 instances), and C (7 instances)
- $\text{need} = \text{max} - \text{allocation}$

	allocation	max	need
	A B C	A B C	A B C
P_0	0 1 0	7 5 3	7 4 3
P_1	2 0 0	3 2 2	1 2 2
P_2	3 0 2	9 0 2	6 0 0
P_3	2 1 1	2 2 2	0 1 1
P_4	0 0 2	4 3 3	4 3 1

available
A B C
3 3 2

Banker's Algorithm: Example

- System state:
 - **5 processes** P_0 through P_4
 - **3 resource types**: A (10 instances), B (5 instances), and C (7 instances)
- $\text{need} = \text{max} - \text{allocation}$

	allocation	max	need
	A B C	A B C	A B C
P_0	0 1 0	7 5 3	7 4 3
P_1	2 0 0	3 2 2	1 2 2
P_2	3 0 2	9 0 2	6 0 0
P_3	2 1 1	2 2 2	0 1 1
P_4	0 0 2	4 3 3	4 3 1

available
A B C
3 3 2

- First one can be either P_1 or P_3

Banker's Algorithm: Example

- What's the safe sequence

	allocation	max	need
	A B C	A B C	A B C
P ₀	0 1 0	7 5 3	7 4 3
P ₁	2 0 0	3 2 2	1 2 2
P ₂	3 0 2	9 0 2	6 0 0
P ₃	2 1 1	2 2 2	0 1 1
P ₄	0 0 2	4 3 3	4 3 1

available
A B C
3 3 2

- finish[1] = true, needed[1] < work -> work = work + allocation = [5 3 2]
- finish[3] = true, needed[3] < work -> work = work + allocation = [7 4 3]
- finish[4] = true, needed[4] < work -> work = work + allocation = [7 4 5]
- finish[2] = true, needed[2] < work -> work = work + allocation = [10 4 7]
- finish[0] = true, needed[0] < work -> work = work + allocation = [10 5 7]

Banker's Algorithm: Example

- P1 requires and gets allocated 1 0 2 more

	allocation	max	need
	A B C	A B C	A B C
P ₀	0 1 0	7 5 3	7 4 3
P ₁	3 0 2	4 2 4	1 2 2
P ₂	3 0 2	9 0 2	6 0 0
P ₃	2 1 1	2 2 2	0 1 1
P ₄	0 0 2	4 3 3	4 3 1


available
A B C
2 3 0

- Check whether it is in safe state?
 - We cannot find a process that the $\text{need}[i] < \text{work}[i]$

Deadlock Detection

- Allow system to enter deadlock state, but detect and recover from it
- Detection algorithm and recovery scheme

09: Main Memory




Logical vs. Physical Address Space

- The concept of a logical address space that is bound to a separate physical address space is central to proper memory management
 - Logical address – generated by the CPU; also referred to as virtual address
 - Physical address – address seen by the memory unit
- Logical address space is the set of all logical addresses generated by a program
- Physical address space is the set of all physical addresses generated by a program

Memory-Management Unit (MMU)

- Hardware device that at run time maps virtual to physical address



Contiguous Allocation

- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one early method
- Main memory usually into two **partitions**:
 - Resident operating system, usually held in low memory with interrupt vector
 - User processes then held in high memory
 - Each process contained in single contiguous section of memory

Memory Allocation Strategies

- Fixed partitions
- Variable partitions

Partitioning Strategies – Fixed

- Fixed Partitions – divide memory into equal sized pieces (except for OS)
 - Degree of multiprogramming = number of partitions
 - Simple policy to implement
 - All processes must fit into partition space
 - Find any free partition and load process
- Problem – **Internal Fragmentation**
 - Unused memory in partition not available to other processes

Question:

What is the “right” partition size?

Partitioning Strategies – Variable

- Memory is dynamically divided into partitions based on process needs
 - More complex management problem
 - Need data structures to track free and used memory
 - New process allocated memory from hole large enough to fit it
- Problem – External Fragmentation
 - Unused memory between partitions too small to be used by any processes

Memory Allocation

- How to satisfy a request of size n from a list of free memory blocks?
 - **first-fit**: allocate from the first block that is big enough
 - **best-fit**: allocate from the smallest block that is big enough
 - must search entire list, unless ordered by size
 - produces the smallest leftover hole
 - **worst-fit**: allocate from the largest hole
 - must also search entire list
 - produces the largest leftover hole
- **Fragmentation** is big problem for all three methods
 - first-fit and best-fit usually perform better than worst-fit

Fragmentation

- **External fragmentation**
 - Unusable memory between allocated memory blocks
 - **Total amount** of free memory space **is larger than a request**
 - The request cannot be fulfilled because the free memory is not **contiguous**
 - External fragmentation can be reduced by **compaction**
 - Shuffle memory contents to place all free memory in one large block
 - Program needs to be **relocatable** at runtime
 - Performance overhead, timing to do this operation
 - Another solution: **paging**

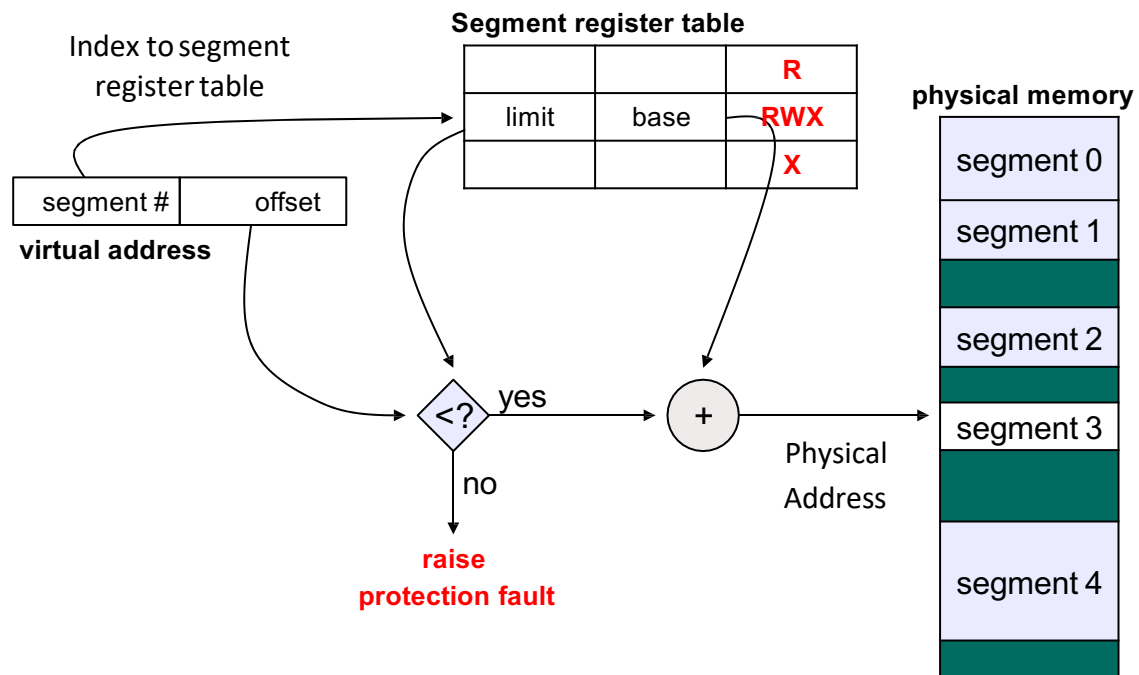
Fragmentation

- **Internal fragmentation**
 - memory allocated may be larger than the requested size
 - this size difference is memory *internal to a partition*, but not being used
 - Example: free space 18464 bytes, request 18462 bytes
- Sophisticated algorithms are designed to avoid fragmentation
 - none of the first-/best-/worst-fit can be considered sophisticated

Segmentation

- Logical address consists of a pair:
 - <segment-number, offset>
- Segment table where each entry has:
 - Base: starting physical address
 - Limit: length of segment

Segment Lookup



Address Binding

- Addresses are represented in different ways at different stages of a program's life
 - **source code** addresses are usually **symbolic** (e.g., variable name)
 - **compiler** binds symbols to **relocatable addresses**
 - e.g., “14 bytes from beginning of this module”
 - **linker** (or loader) binds relocatable addresses to **absolute addresses**
 - e.g., 0x0e74014
 - Each binding maps one address space to another

Binding of Instructions and Data to Memory

- Address binding of instructions and data to memory addresses can happen at three different stages
 - **Compile time:** If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes
 - **Load time:** Must generate relocatable code if memory location is not known at compile time
 - **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - Need **hardware support** for address maps (e.g., base and limit registers)

Paging

- Physical address space of a process can be **noncontiguous**; process is allocated physical memory whenever the latter is available
 - Avoids **external fragmentation** -> **avoid for compacting**
 - Avoids problem of varying sized memory chunks
- Basic methods
 - Divide physical memory into fixed-sized blocks called **frames**
 - Size is power of 2, between 512 bytes and 16 Mbytes
 - Divide logical memory into blocks of same size called **pages**
 - Keep track of all free frames
 - To run a program of size **N** pages, need to find **N** free frames and load program
 - Set up a **page table** to translate logical to physical addresses
 - Backing store likewise split into pages
 - **Still have Internal fragmentation, where?**


Paging: Address Translation

- A logical address is divided into:
 - **page number** (p)
 - used as an index into a page table
 - page table entry contains the corresponding **physical** frame number
 - **page offset** (d)
 - offset within the page/frame
 - combined with frame number to get the physical address

page number	page offset
p	d
$m - n \text{ bits}$	$n \text{ bits}$

m bit logical address space, n bit page size

Paging Example



$m = 4$ and $n = 2$ 32-byte memory and 4-byte pages

TLB

- TLB and context switch
 - Each process has its own page table
 - Switching process needs to switch page table
 - **TLB must be consistent with page table**
 - TLB entries are from page table of current process
 - Option I: flush TLB at every context switch, **or**,
 - Option II: tag TLB entries with **address-space identifier** (ASID) that uniquely identifies a process
 - Some TLB entries can be **shared** by processes, and fixed in the TLB
 - E.G., TLB entries for the kernel
- TLB and operating system
 - MIPS: OS should deal with TLB miss exception
 - X86: TLB miss is handled by hardware

Effective Access Time

- **Hit ratio** – percentage of times that a page number is found in the TLB
- An 80% hit ratio means that we find the desired page number in the TLB 80% of the time.
- Suppose that 10 nanoseconds to access memory.
 - If we find the desired page in TLB then a mapped-memory access take 10 ns
 - Otherwise we need **two memory access** so it is 20 ns: page table + memory access
- Effective Access Time (EAT)
- $$\text{EAT} = 0.80 \times 10 + 0.20 \times 20 = 12 \text{ nanoseconds}$$
- Implying 20% slowdown in access time
- What if TLB access time is non-zero
 - For example, TLB access time is 2 ns, memory access is 10 ns, hit ratio is p
 - With TLB: total time = $p \times (2+10) + (1-p) \times (2+20)$
 - Without TLB: 20


Valid-Invalid Bit

- Each page table entry has a valid–invalid (present) bit
 - \underline{V} \Rightarrow in memory (memory is resident), \underline{I} \Rightarrow not-in-memory
 - initially, valid–invalid bit is set to \underline{I} on all entries
 - during address translation, if the entry is invalid, it will trigger a **page fault**
- Example of a page table snapshot:

Frame #	v/i bit
	v
	v
	v
	v
	i
....	
	i
	i


page table

Page Table (Some Pages Are Not in Memory)

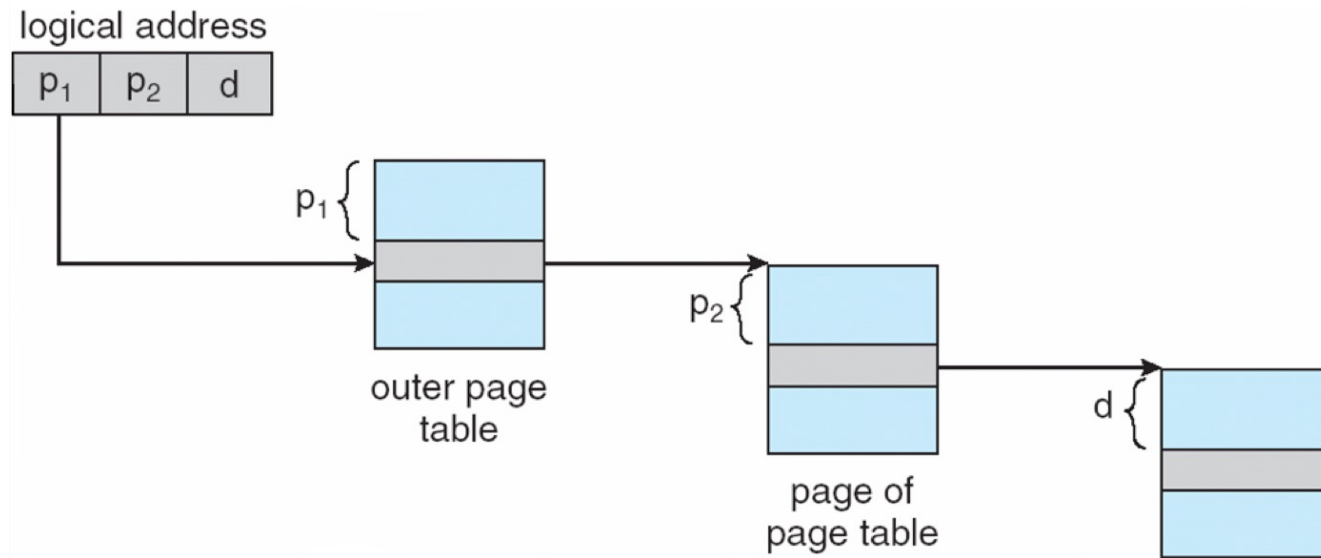


Two-Level Paging

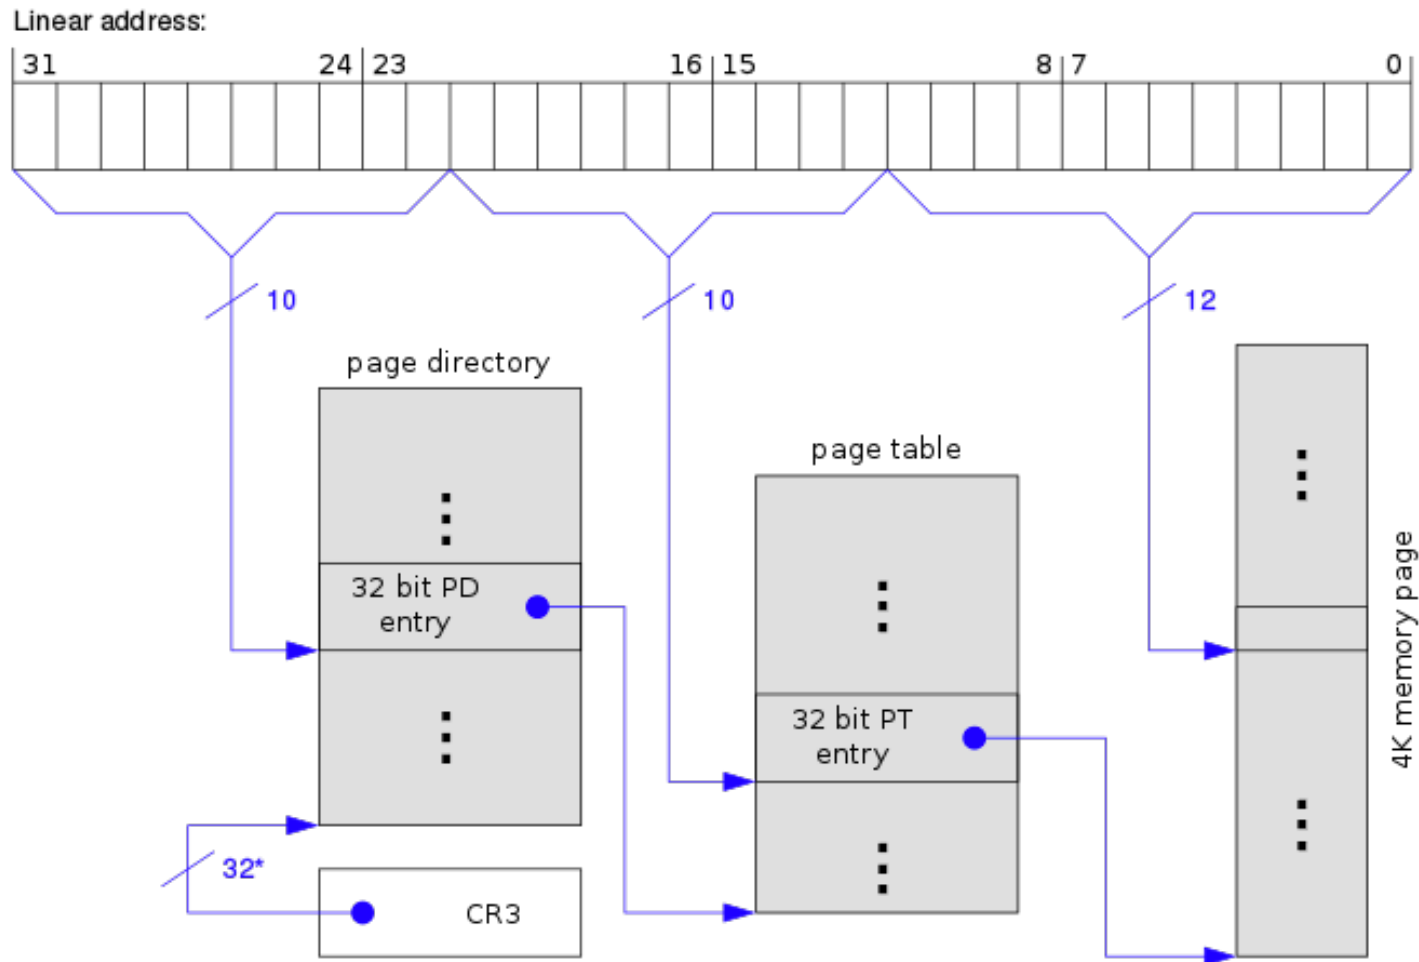
- A logical address is divided into:
 - A **page directory number** (first level page table)
 - A **page table number** (2nd level page table)
 - A **page offset**
- Example: 2-level paging in 32-bit intel cpus
 - 32-bit address space, 4KB page size
 - 10-bit page directory number, 10-bit page table number
 - Each page table entry is 4 bytes, one frame contains 1024 entries (2^{10})



Address-Translation Scheme



Page Table in Linux



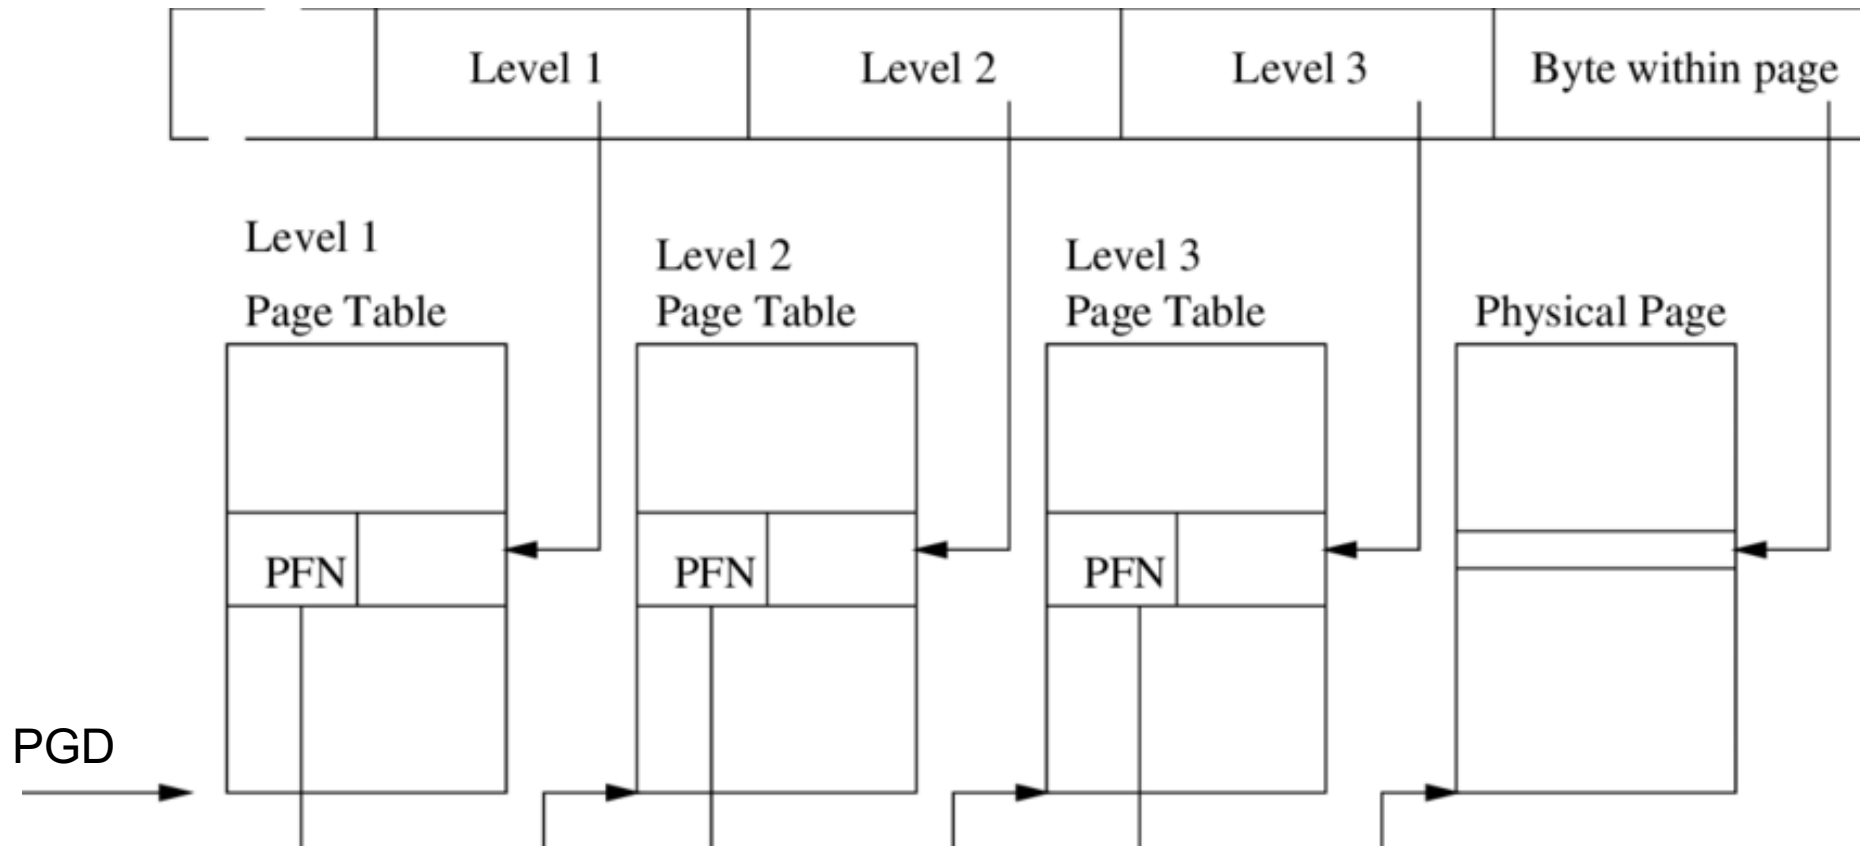
*) 32 bits aligned to a 4-KByte boundary

64-bit Logical Address Space

- 64-bit logical address space requires more levels of paging
 - two-level paging is not sufficient for 64-bit logical address space
 - if page size is 4 KB (2^{12}), outer page table has 2^{42} entries, inner page tables have 2^{10} 4-byte entries
 - one solution is to add more levels of page tables
 - e.g., three levels of paging: 1st level page table is 2^{32} bytes in size
 - and possibly 4 memory accesses to get to one physical memory location
- usually **not support full 64-bit virtual address space**
 - AMD-64 supports 48-bit
 - ARM64 supports 39-bit, 48-bit

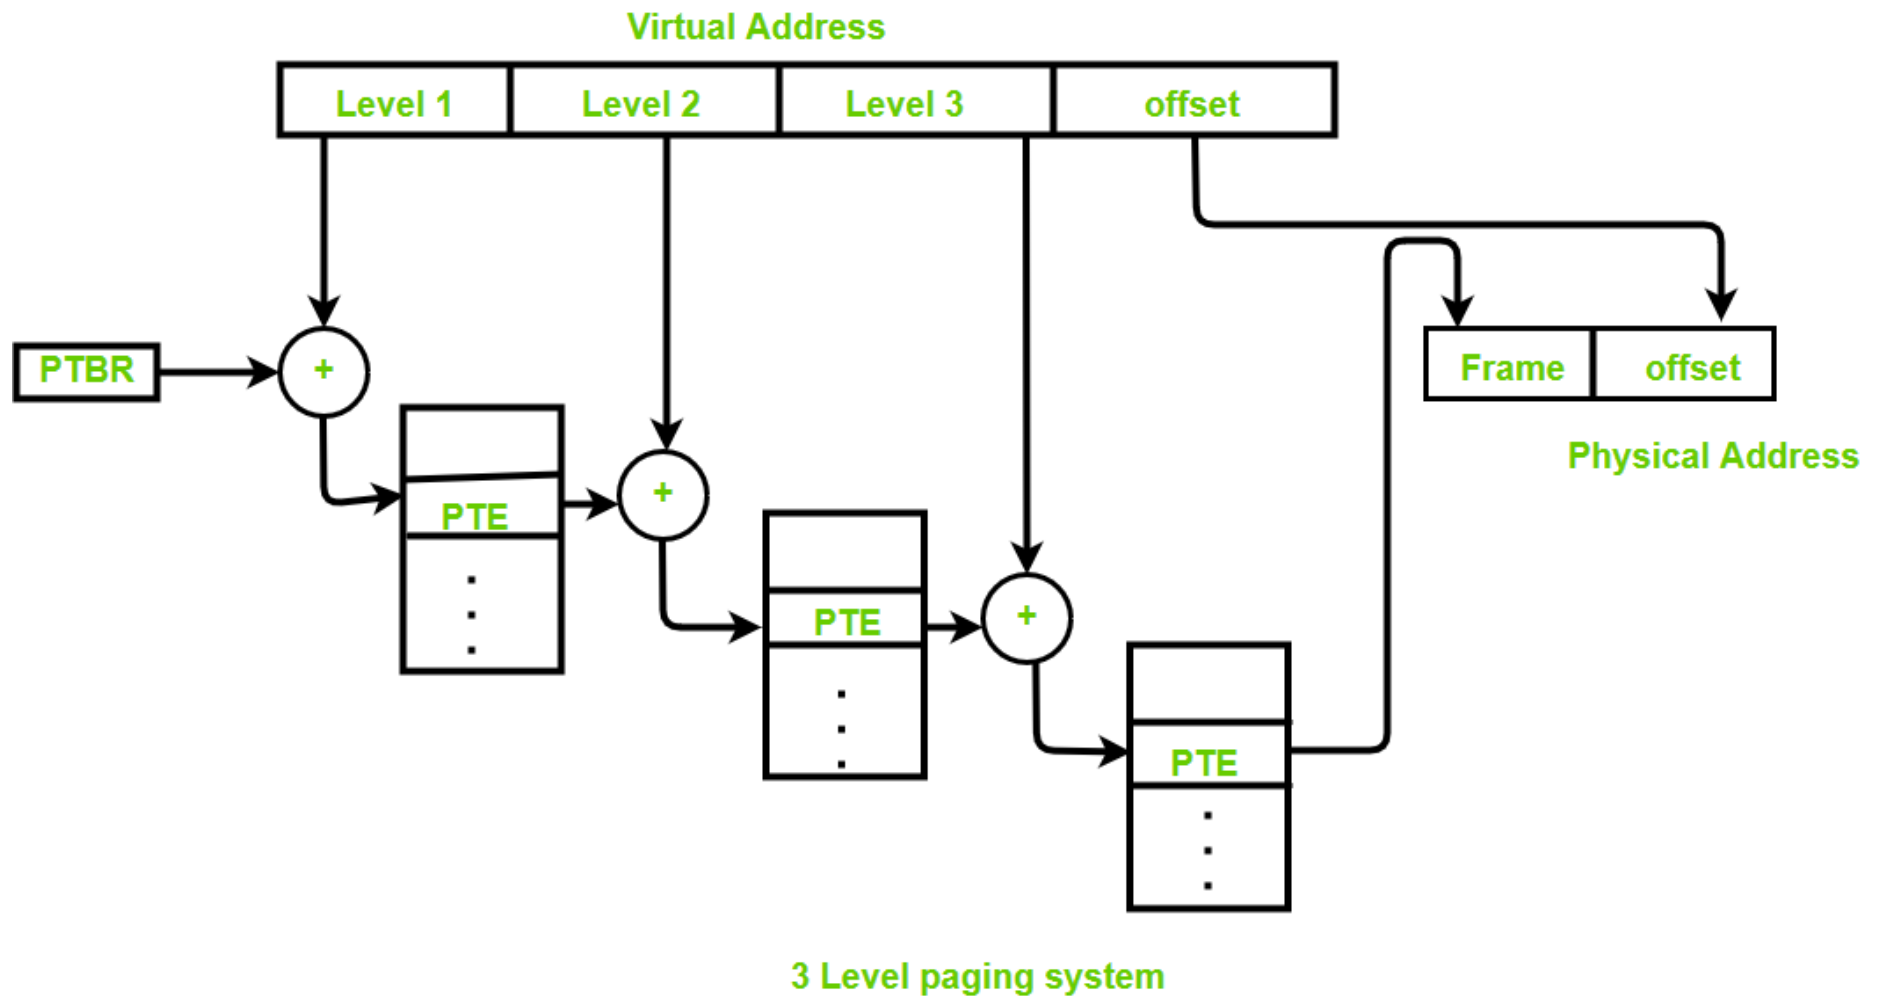
64-bit Logical Address Space

- ARM64: 39 bits = 9+9+9+12



64-bit Logical Address Space

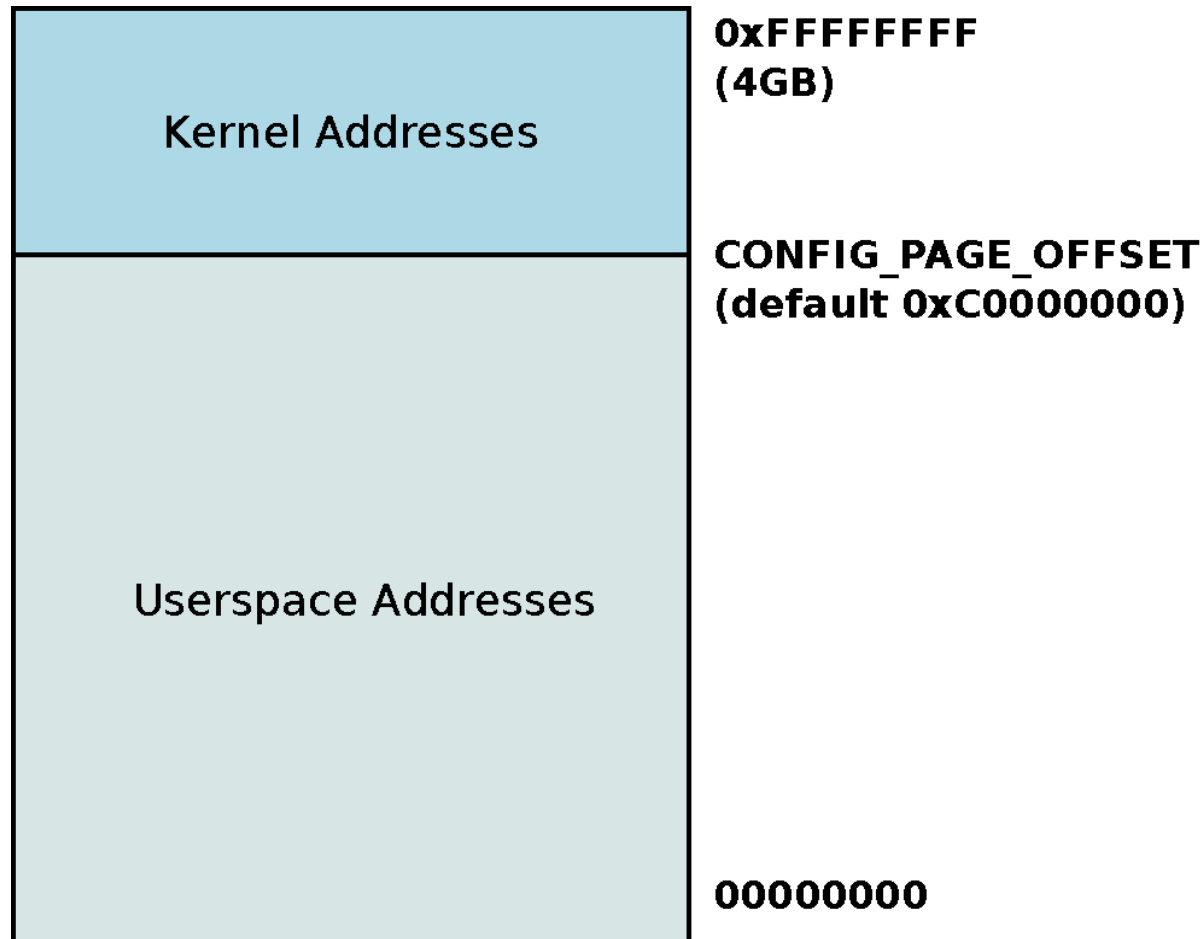
- ARM64: 39 bits = 9+9+9+12



Memory Protection

- Accomplished by protection bits with each frame
- Each page table entry has a **present** (aka. valid) bit
 - present: the page has a valid physical frame, thus can be accessed
- Each page table entry contains some protection bits
 - **kernel/user, read/write, execution?, kernel-execution?**
 - why do we need them?
- Any violations of memory protection result in a trap to the kernel

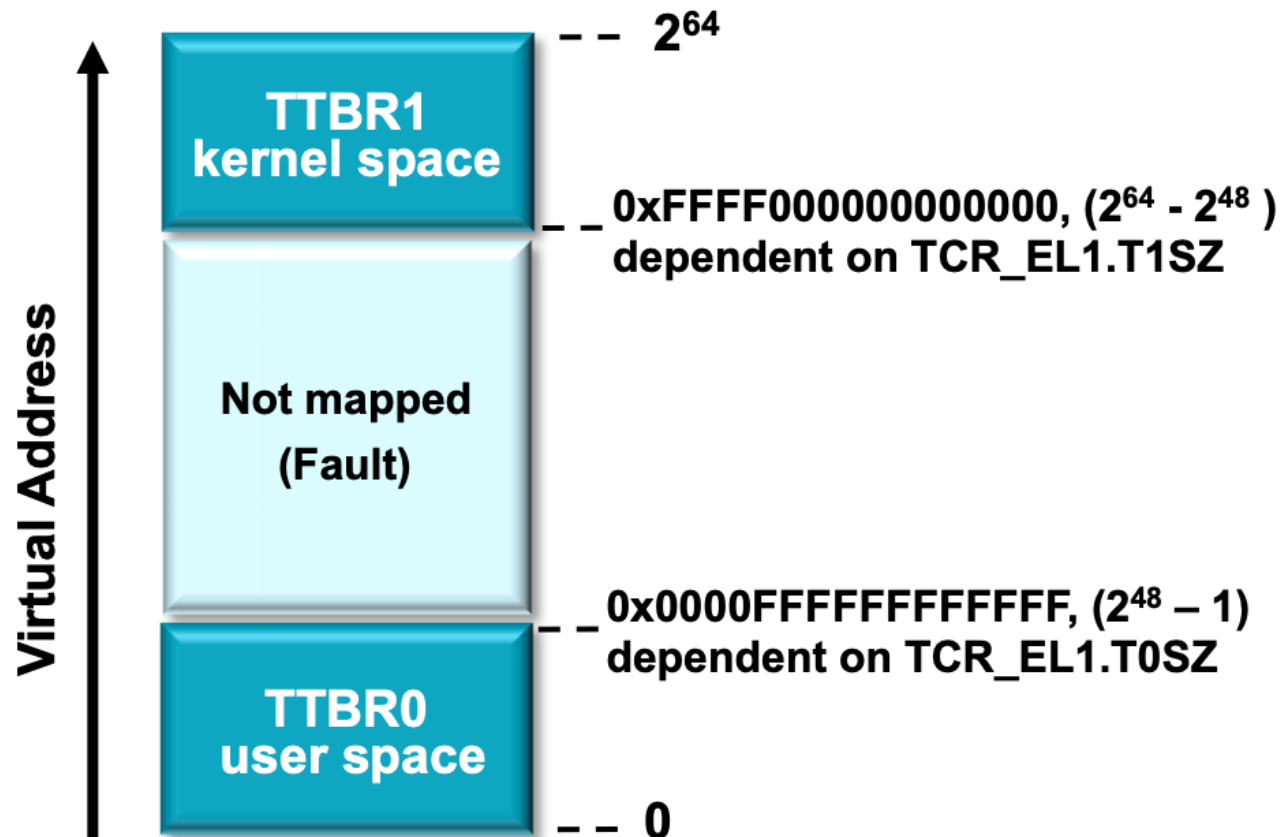
Virtual Addresses 32-bit – Linux



- By default, the kernel uses the top 1GB of virtual address space.
- Each userspace processes get the lower 3GB of virtual address space.

Virtual Addresses 64-bit

- ARM 64
 - 48-bit page table
 - 4-level: 9+9+9+9+12

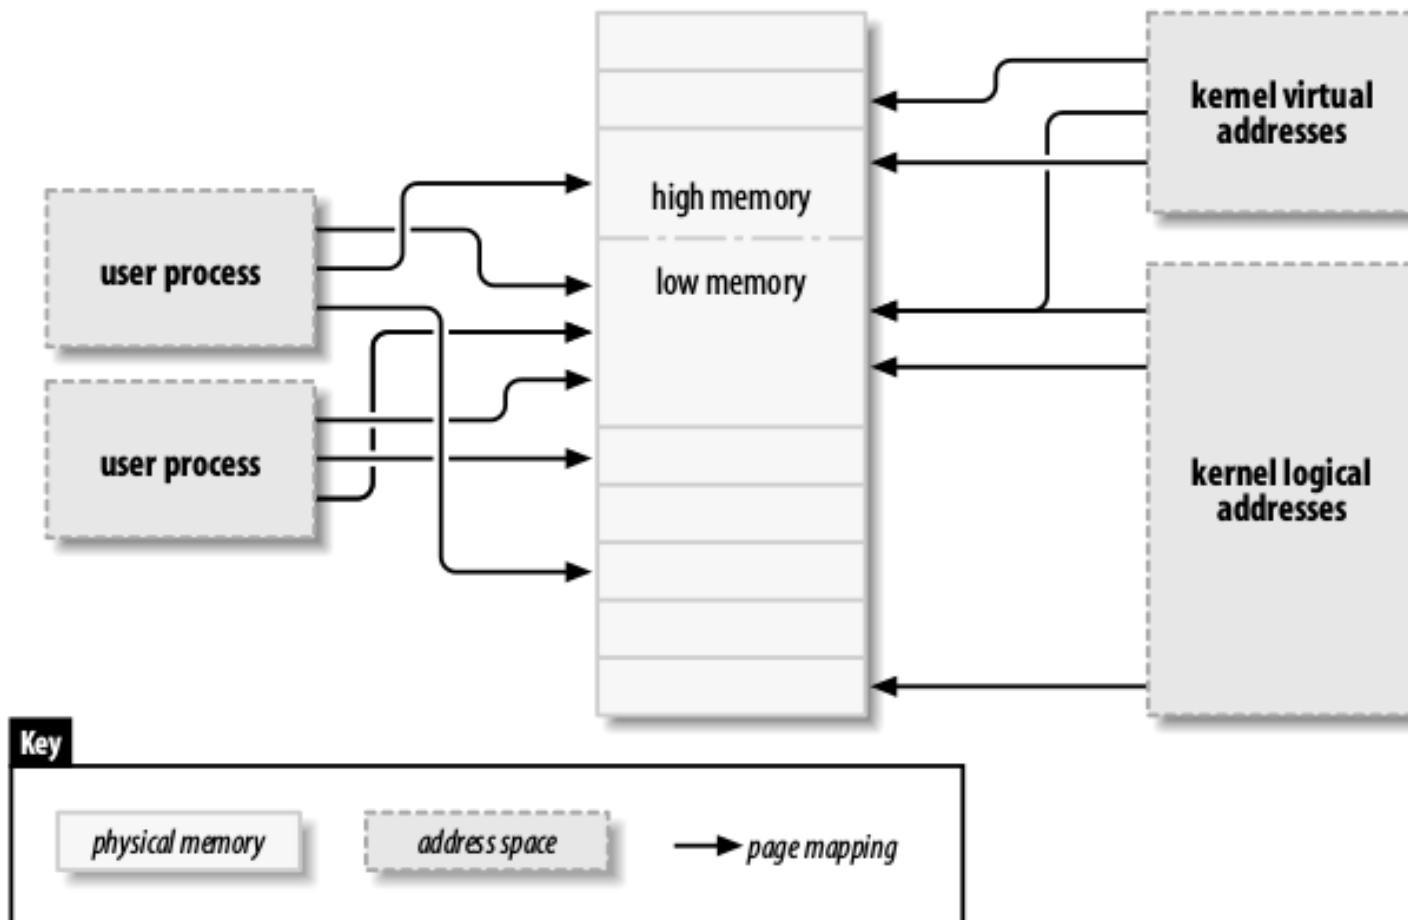


Virtual Addresses 64-bit


- ARM 64
 - 39-bit page table, 3-level: 9+9+9+12
 - 39-bit virtual address for both user and kernel
 - 0000000000000000–0000007fffffffff (512GB): user
 - [architectural gap]
 - ffffffff8000000000–ffffffbfffffff (∼240MB): vmalloc
 - ffffffffbfffffff0000–ffffffbfffffff (64KB): [guard]
 - ffffffffb000000000–ffffffbfffffffff (8GB): vmemmap
 - ffffffffb000000000–ffffffbffffffffff (∼8GB): [guard]
 - ffffffffbffc000000–ffffffbfffffffff (64MB): modules
 - ffffffffc000000000–ffffffffffffffff (256GB): mapped RAM
 - 4KB page configuration
 - 3 levels of page tables (pgtable-nopud.h)
 - Linear mapping using 4KB, 2MB or 1GB blocks
 - AArch32 (compat) supported

Virtual Addresses - Linux

- User virtual address
- Kernel virtual address
- Kernel logical address



10: Virtual Memory



Demand Paging Background

- Code needs to be in memory to execute, but entire program **rarely** needed or used at the same time
 - **unused code**: error handling code, unusual routines
 - **unused data**: large data structures
- Consider ability to execute **partially-loaded program**
 - program no longer constrained by limits of physical memory
 - programs could be larger than physical memory

Demand Paging

- **Demand paging** brings a **page** into memory only when it is **demanded**
 - demand means access (read/write)
 - if page is invalid (error) \Rightarrow abort the operation
 - if page is valid but not in memory \Rightarrow bring it to memory
 - Memory here means **physical** memory
 - This is called **page fault**
 - via swapping for swapped pages
 - via mapping for new page
 - no unnecessary I/O, less memory needed, slower response, more apps

Page Fault

- First reference to a non-present page will trap to kernel: page fault, the reasons can be
 - **invalid reference** \Rightarrow deliver an exception to the process
 - **valid but not in memory** \Rightarrow swap in
- get an empty physical frame
- swap page into frame via disk operation
- set page table entry to indicate the page is now in memory
- restart the instruction that caused the page fault


What Happens if There is no Free Frame?

- Used up by process pages
- Also in demand from the kernel, I/O buffers, etc
- How much to allocate to each?
- **Page replacement** – find some page in memory, but not really in use, page it out
 - Algorithm – terminate? swap out? replace the page?
 - Performance – want an algorithm which will result in minimum number of page faults
- Same page may be brought into memory several times


Page Replacement

- Memory is an important resource, system may run out of memory
- To prevent out-of-memory, swap out some pages
 - Page replacement usually is a part of the page fault handler
 - Policies to select victim page require careful design
 - Need to reduce overhead and avoid thrashing
 - Use modified (dirty) bit to reduce number of pages to swap out
 - Only modified pages are written to disk
 - Select some processes to kill (last resort)
- Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory

Need For Page Replacement



Page Replacement




Page Replacement Algorithms

- Page-replacement algorithm should have lowest page-fault rate on both first access and re-access
 - **FIFO, optimal, LRU, LFU, MFU...**
- To evaluate a page replacement algorithm:
 - Run it on a particular string of memory references (reference string)
 - String is just page numbers, not full addresses
 - Compute the number of page faults on that string
 - Repeated access to the same page does not cause a page fault
 - In all our examples, the reference string is
7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1

First-In-First-Out (FIFO)


- **FIFO**: replace the first page loaded
 - similar to sliding a window of n in the reference string
 - our reference string will cause 15 page faults with 3 frames
 - how about reference string of 1,2,3,4,1,2,5,1,2,3,4,5 /w 3 or 4 frames?
- For FIFO, adding **more frames** can cause **more page faults**!
- **Belady's Anomaly**



15 page faults


Optimal Algorithm

- **Optimal** : replace page that will not be used for the longest time
 - 9 page fault is optimal for the example on the next slide
- **How do you know which page will not be used for the longest time?**
 - **can't read the future**
 - used for measuring how well your algorithm performs



Least Recently Used (LRU)

- **LRU** replaces pages that **have not been used for the longest time**
 - associate time of last use with each page, select pages w/ oldest timestamp
 - generally good algorithm and frequently used
 - 12 faults for our example, better than FIFO but worse than OPT
- LRU and OPT do **NOT** have Belady's Anomaly



Thrashing

- If a process doesn't have "enough" pages, page-fault rate may be high
 - page fault to get page, replace some existing frame
 - but quickly need replaced frame back
 - this leads to:

low CPU utilization ➡


kernel thinks it needs to increase the degree of

multiprogramming to maximize CPU utilization ➡

another process added to the system

- **Thrashing:** a process is busy swapping pages in and out

Thrashing



Demand Paging and Thrashing


- Why does demand paging work?
 - process memory access has **high locality**
 - process migrates from one locality to another, localities may overlap
- Why does thrashing occur?
 - total size of locality > total memory size

Working-Set Model

- **Working-set window(Δ):** a fixed number of page references
 - if Δ too small \Rightarrow will not include entire locality
 - if Δ too large \Rightarrow will include several localities
 - if $\Delta = \infty \Rightarrow$ will include entire program
- **Working set** of process p_i (WSS_i): total number of pages referenced in the most recent Δ (varies in time)
- **Total working sets:** $D = \sum WSS_i$
 - approximation of total locality
 - if $D > m \Rightarrow$ possibility of thrashing
 - to avoid thrashing: if $D > m$, suspend or swap out some processes

Working-Set Model

- Working-set window $\Delta = 10$



Page table quiz

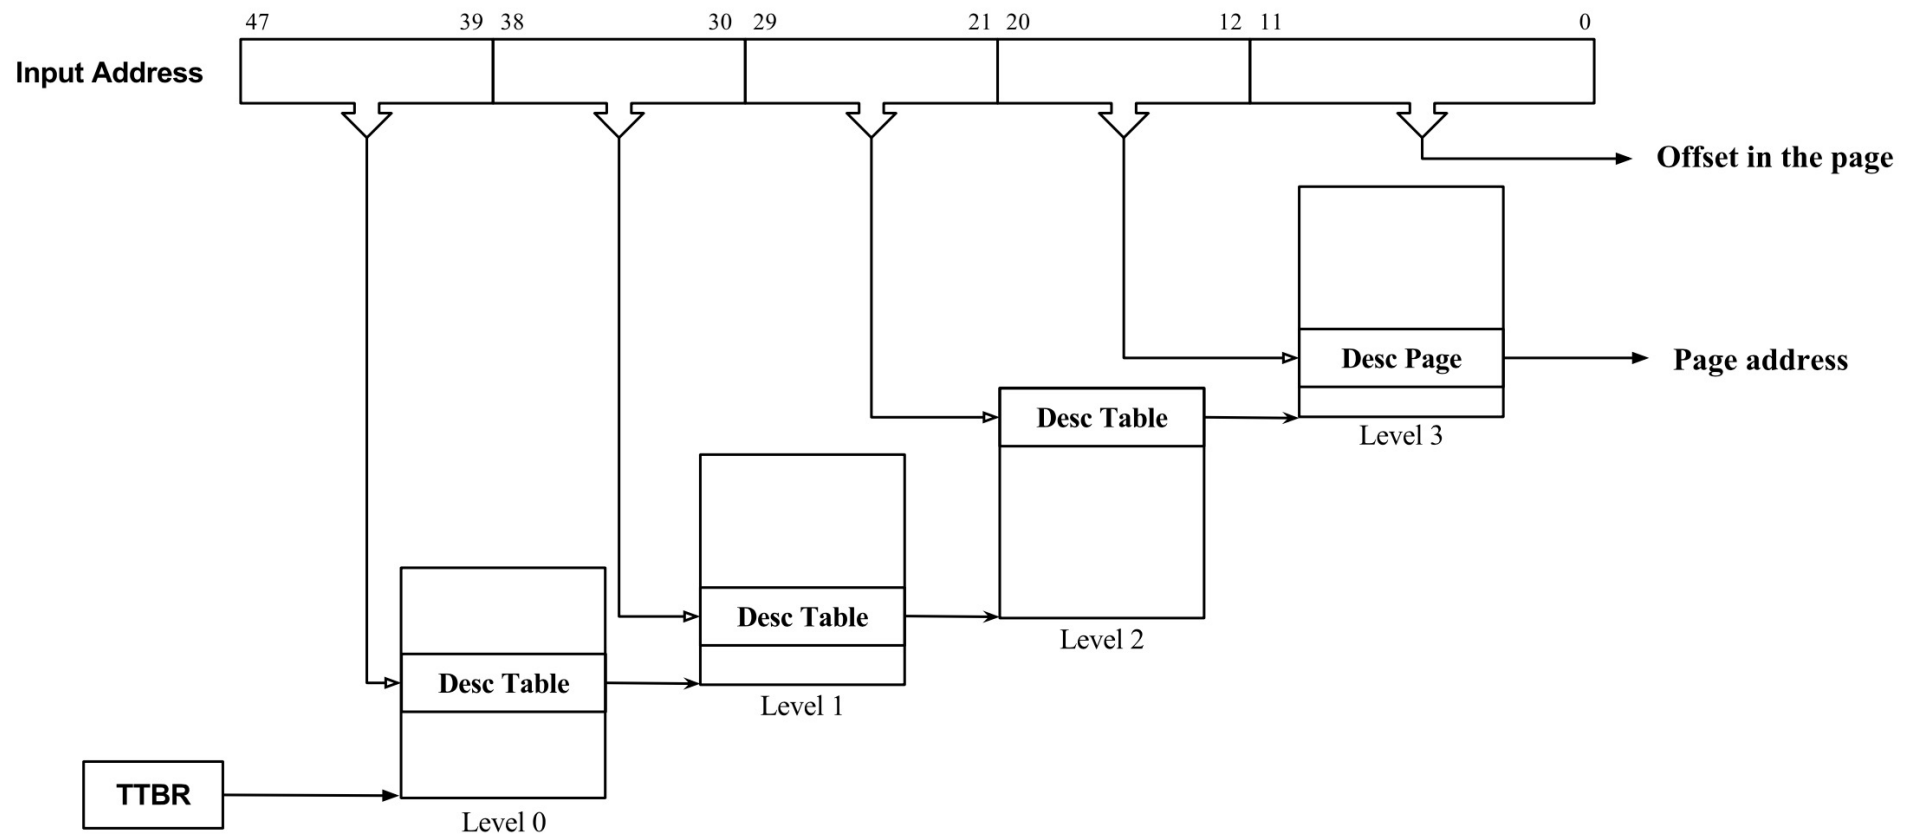
- 1) In 32-bit architecture, for 1-level page table, how large is the whole page table?
- 2) In 32-bit architecture, for 2-level page table, how large is the whole page table?
 - 1) How large for the 1st level PGT?
 - 2) How large for the 2nd level PGT?
- 3) Why can 2-level PGT save memory?
- 4) 2-level page table walk example
 - 1) Page table base register holds 0x0061 9000
 - 2) Virtual address is 0xf201 5202
 - 3) Page table base register holds 0x1051 4000
 - 4) Virtual address is 0x2190 7010

Page table quiz

- How about page size is 64KB
 - What is the virtual address format for 32-bit?
 - What is the virtual address format for 64-bit?
 - 39-bit VA
 - 48-bit VA

Virtual address format

- 48-bit VA with 4KB page



Takeaway

- The whole slides